

The MIPS instruction set architecture


The MIPS has a 32 bit architecture, with 32 bit instructions, a 32 bit data word, and 32 bit addresses.

It has 32 addressable internal registers requiring a 5 bit register address. Register 0 always has the the constant value 0.


Addresses are for individual bytes (8 bits) but instructions must have addresses which are a multiple of 4. This is usually stated as “instructions must be word aligned in memory.”

There are three basic instruction types with the following formats:


R-type (register)



I-type (immediate)



J-type (jump)



All op codes are 6 bits.

All register addresses are 5 bits.

R-type (register)

31	26	25	21	20	16	15	11	10	6	5	0		
op				rs		rt		rd		shamt		funct	

The R-type instructions are 3 operand arithmetic and logic instructions, where the operands are contained in the registers indicated by **rs**, **rt**, and **rd**.

For all R-type instructions, the **op** field is 000000.

The **funct** field selects the particular type of operation for R-type operations.

The **shamt** field determines the number of bits to be shifted (0 to 31).


These instructions perform the following:

$$R[rd] \leftarrow R[rs] \text{ op } R[rt]$$

Following are examples of R-type instructions:

Instruction	Example	Meaning
add	add \$s1, \$s2, \$s3	\$s1 = \$s2 + \$s3
add unsigned	addu \$s1, \$s2, \$s3	\$s1 = \$s2 + \$s3
subtract	sub \$s1, \$s2, \$s3	\$s1 = \$s2 - \$s3
subtract unsigned	subu \$s1, \$s2, \$s3	\$s1 = \$s2 - \$s3
and	and \$s1, \$s2, \$s3	\$s1 = \$s2 & \$s3
or	or \$s1, \$s2, \$s3	\$s1 = \$s2 \$s3

I-type (immediate)



The 16 bit **immediate** field contains a data constant for an arithmetic or logical operation, or an address offset for a **branch** instruction. This type of branch is called a *relative branch*.

Following are examples of I-type instructions of type:

$$R[rt] \leftarrow R[rs] \text{ op } imm$$

Instruction	Example	Meaning
add	<code>addi \$s1, \$s2, imm</code>	$\$s1 = \$s2 + imm$
add unsigned	<code>addiu \$s1, \$s2, imm</code>	$\$s1 = \$s2 + imm$
subtract	<code>subi \$s1, \$s2, imm</code>	$\$s1 = \$s2 - imm$
and	<code>andi \$s1, \$s2, imm</code>	$\$s1 = \$s2 \& imm$


Another I-type instruction is the **branch** instruction.

Examples of this are:

Instruction	Example	Meaning
branch on equal	<code>beq \$s1, \$s2, imm</code>	if $\$s1 == \$s2$ go to $PC + 4 + (4 \times imm)$
branch on not equal	<code>bne \$s1, \$s2, imm</code>	if $\$s1 != \$s2$ go to $PC + 4 + (4 \times imm)$

Why is the **imm** field multiplied by 4 here?

J-type (jump)



The J-type instructions are all **jump** instructions.

The two we will discuss are the following:

Instruction	Example	Meaning
jump	j target	go to address $4 \times \mathbf{target} : \text{PC}[28:31]$
jump and link	jal target	$\$31 = \text{PC} + 4;$ go to address $4 \times \mathbf{target} : \text{PC}[28:31]$

Why is the **PC** incremented by 4?

Why is the **target** field multiplied by 4?

Recall that the MIPS processor addresses data at the *byte* level, but instructions are addressed at the **word** level.

Moreover, all instructions must be aligned on a word boundary (an integer multiple of 4 bytes).

Therefore, the next instruction is 4 byte addresses from the current instruction.

Since jumps must have an instruction as target, shifting the target address by 2 bits (which is the same as multiplying by 4) allows the instruction to specify larger jumps.

Note that the **jump** instruction cannot span (jump across) all of memory.

There are a few more interesting instructions, for comparison, and memory access:

R-type instructions:

Instruction	Example	Meaning
set less than	<code>slt \$s1, \$s2, \$s3</code>	if ($\$s2 < \$s3$), $\$s1=1$; else $\$s1=0$
jump register	<code>jr \$ra</code>	go to $\$ra$

set less than also has an unsigned form.

jump register is typically used to return from a subprogram.

I-type instructions:

Instruction	Example	Meaning
set less than immediate	<code>slti \$s1, \$s2, imm</code>	if ($\$s2 < \text{imm}$), $\$s1=1$; else $\$s1=0$
load word	<code>lw \$s1, imm(\$s2)</code>	$\$s1 = \text{Memory}[\$s2 + \text{imm}]$
store word	<code>sw \$s1, imm(\$s2)</code>	$\text{Memory}[\$s2 + \text{imm}] = \$s1$

load word and **store word** are the only instructions that access memory directly.

Because data must be explicitly loaded before it is operated on, and explicitly stored afterwards, the MIPS is said to be a *load/store* architecture.

This is often considered to be an essential feature of a reduced instruction set architecture (RISC).

The MIPS assembly language

The previous diagrams showed examples of code in a general form which is commonly used as a simple kind of language for a processor — a language in which each line in the code corresponds to a single instruction in the language understood by the machine.

For example,

add \$1, \$2, \$3

means take add together the contents of registers **\$2** and **\$3** and store the result in register **\$1**.

We call this type of language an *assembly language*.

The language of the machine itself, called the *machine language*, consists only of 0's and 1's — a binary code.

The machine language instruction corresponding to the previous instruction (with the different fields identified) is:

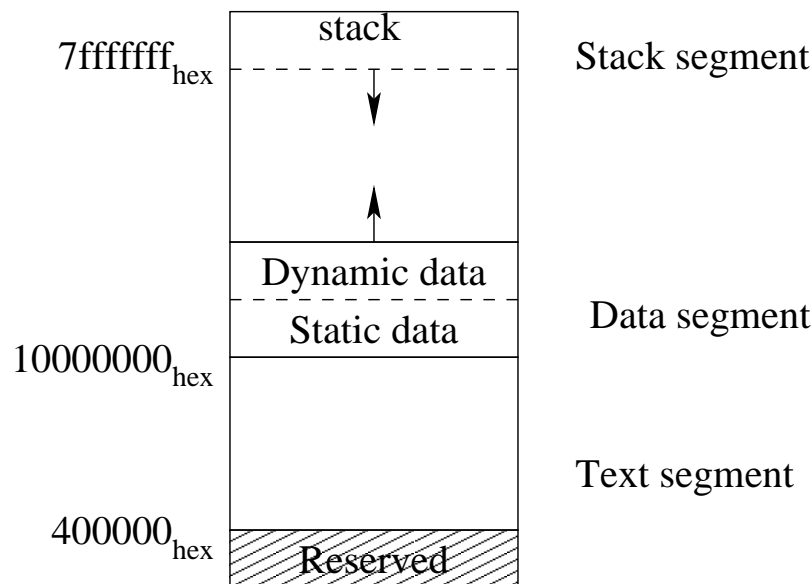
31	26	25	21	20	16	15	11	10	6	5	0
000000	00010	00011	00001	00000	100000						
op	rs	rt	rd	shamt	funct						

There are usually programs, called *assemblers*, to translate the more human readable assembly code to machine language.

MIPS memory usage

MIPS systems typically divided memory into three parts, called *segments*.

These segments are the *text segment* which contains the program's instructions, the *data segment*, which contains the program's data, and the *stack segment* which contains the return addresses for function calls, and also contains register values which are to be saved and restored. It may also contain local variables.



The *data segment* is divided into 2 parts, the lower part for static data (with size known at compile time) and the upper part, which can grow, upward, for dynamic data structures.

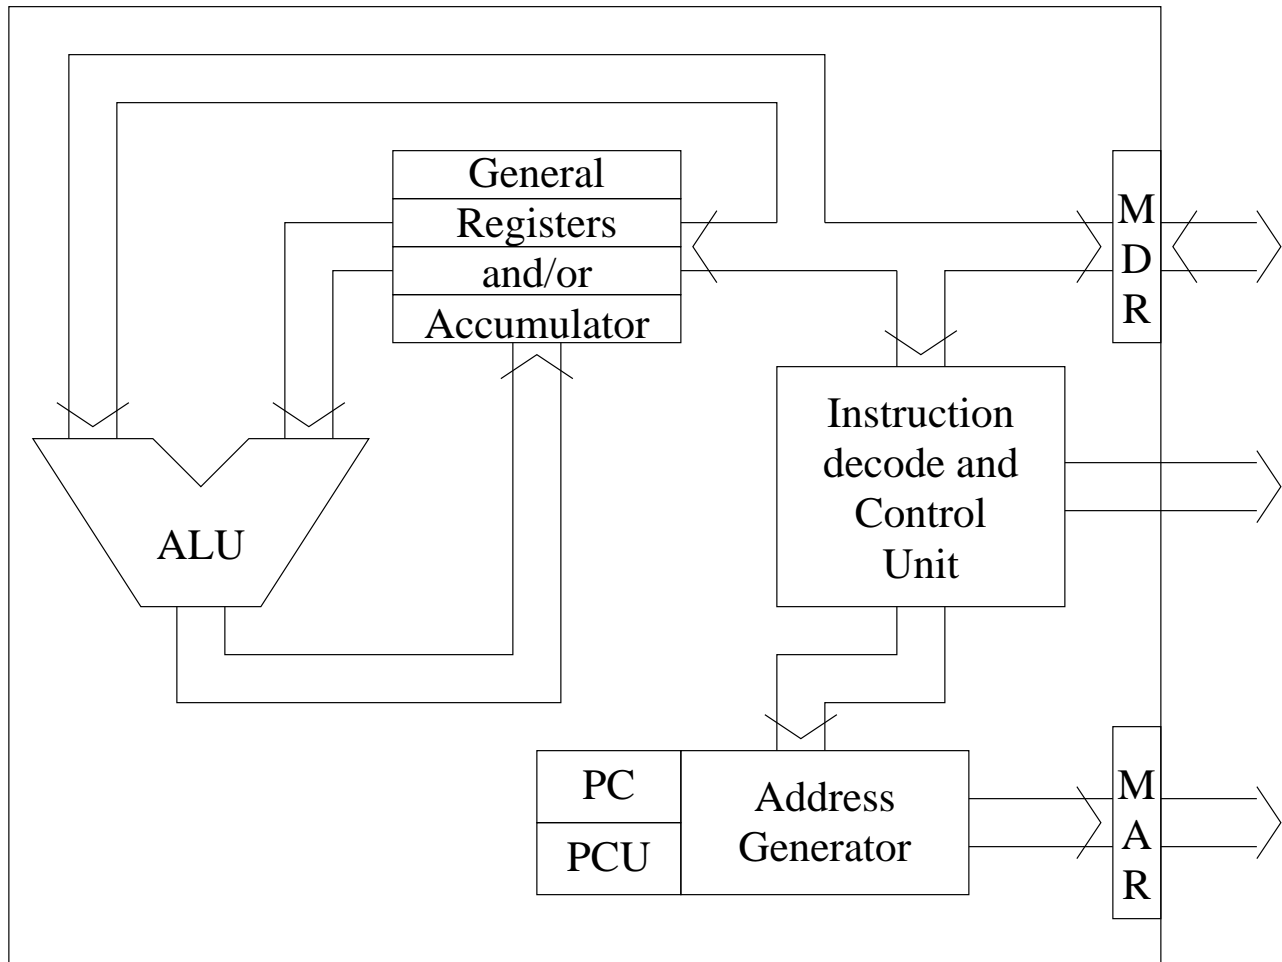
The *stack segment* varies in size during the execution of a program, as functions are called and returned from.

It starts at the top of memory and grows down.

MIPS register names and conventions about their use

Register Name	Number	Usage
zero	0	Constant 0
at	1	Reserved for assembler
v0	2	Expression evaluation and
v1	3	results of a function
a0	4	Argument 1
a1	5	Argument 2
a2	6	Argument 3
a3	7	Argument 4
t0	8	Temporary (not preserved across call)
t1	9	Temporary (not preserved across call)
t2	10	Temporary (not preserved across call)
t3	11	Temporary (not preserved across call)
t4	12	Temporary (not preserved across call)
t5	13	Temporary (not preserved across call)
t6	14	Temporary (not preserved across call)
t7	15	Temporary (not preserved across call)
s0	16	Saved temporary (preserved across call)
s1	17	Saved temporary (preserved across call)
s2	18	Saved temporary (preserved across call)
s3	19	Saved temporary (preserved across call)
s4	20	Saved temporary (preserved across call)
s5	21	Saved temporary (preserved across call)
s6	22	Saved temporary (preserved across call)
s7	23	Saved temporary (preserved across call)
t8	24	Temporary (not preserved across call)
t9	25	Temporary (not preserved across call)
k0	26	Reserved for OS kernel
k1	27	Reserved for OS kernel
gp	28	Pointer to global area
sp	29	Stack pointer
fp	30	Frame pointer
ra	31	Return address (used by function call)

Presently, we are interested in the CPU only, which we concluded would have a structure similar to the following:



The memory address register (MAR) and memory data register (MDR) are the interface to memory.

The ALU and register file are the core of the data path.

The program control unit (PCU) fetches instructions and data, and handles branches and jumps.


The instruction decode unit (IDU) is the control unit for the processor.

The “building blocks”

We have already designed many of the major components for the processor, or have at least identified how they could be implemented. For example, we have already designed an ALU, a data register, and a register file.

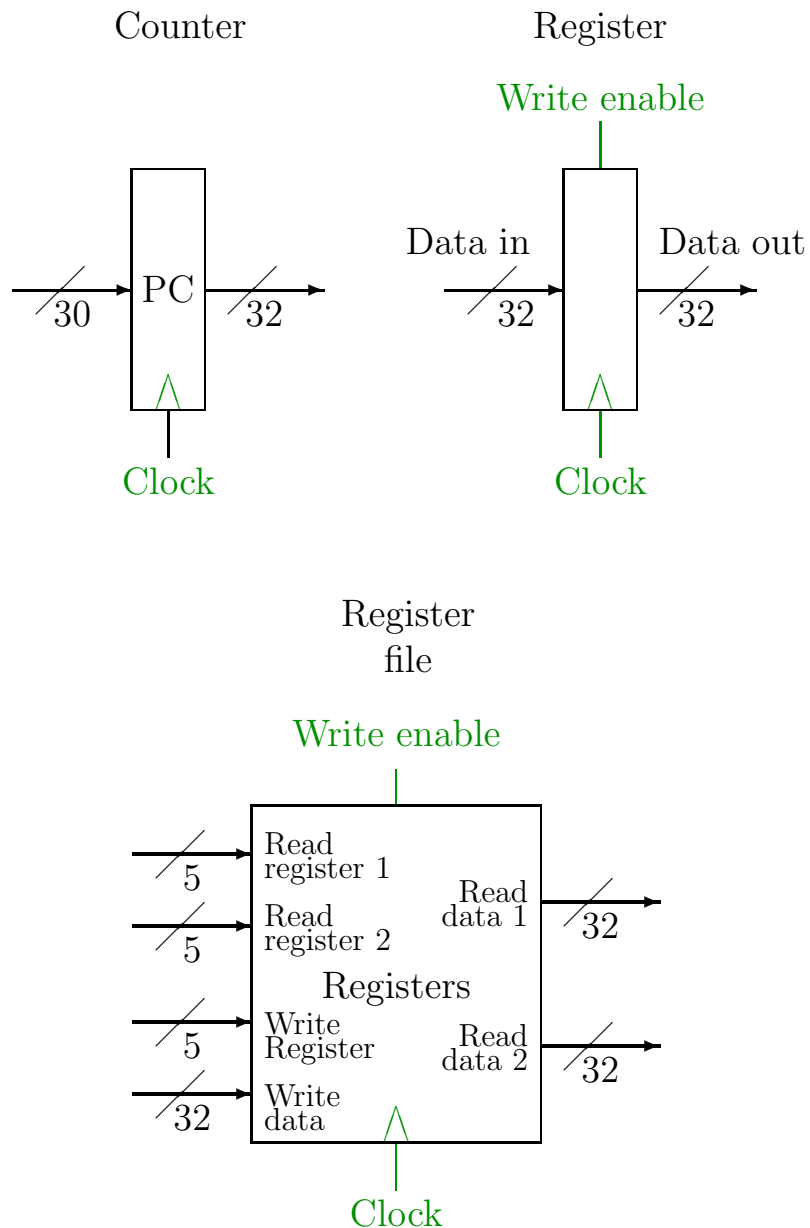
A controller is merely a state machine, and we can implement one using, say, a PLA, after identifying the required states and transitions.

Following are some of the combinational logic components we will use:



Note that the diagram highlights the control signals (**OP** and **S**).

Following are some of the register components we will use:




Note that the registers have a *write enable* input as well as a clock input. This input must be asserted in order for the register to be written.

We have already seen how to construct a register file from simple D registers.

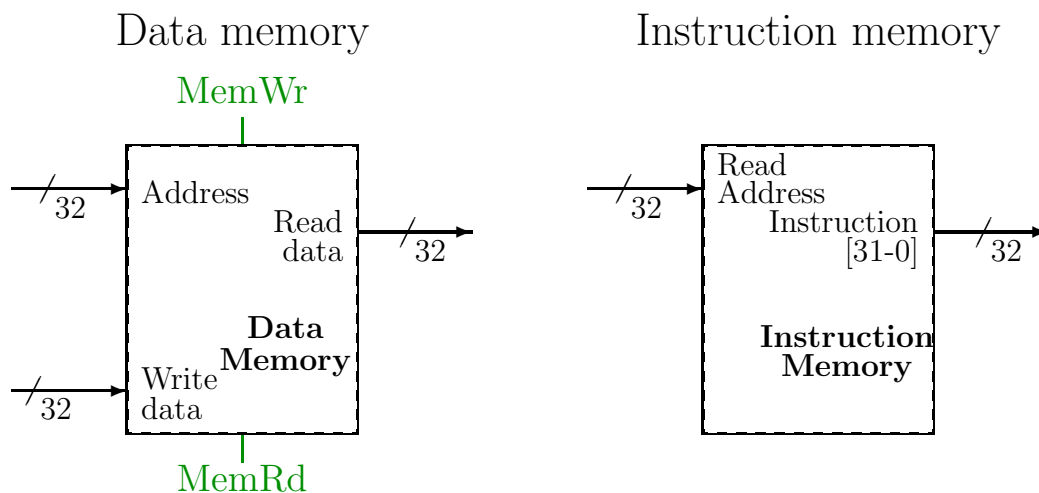
Timing considerations

In a single-cycle implementation of the processor, a single instruction (e.g., **add**) may require that a register be read from and written into in the same clock period. In order to accomplish this, the register file (and other register elements) must be *edge triggered*.

This can be done by using edge triggered elements directly, or by using a *master-slave* arrangement similar to one we saw earlier:




Another observation about a single cycle processor — the memory for instructions must be different from the memory for data, because both must be addressed in the same cycle. Therefore, there must be two memories; one for instructions, and one for data.




The MIPS instruction set:

Following is the MIPS instruction format:


R-type (register)



I-type (immediate)



J-type (jump)



We will develop an implementation of a very basic processor having the instructions:

R-type instructions add, sub, and, or, slt

I-type instructions addi, lw, sw, beq

J-type instructions j

Later, we will add additional instructions.

Steps in designing a processor

- Express the instruction architecture in a Register Transfer Language (RTL)
- From the RTL description of each instruction, determine
 - the required datapath components
 - the datapath interconnections
- Determine the control signals required to enable the datapath elements in the appropriate sequence for each instruction
- Design the control logic required to generate the appropriate control signals at the correct time

A Register Transfer Language description of some operations:

The ADD instruction

add rd, rs, rt

- $\text{mem}[\text{PC}]$ Fetch the instruction from memory
- $\text{R}[\text{rd}] \leftarrow \text{R}[\text{rs}] + \text{R}[\text{rt}]$ Set register **rd** to the value of the sum of the contents of registers **rs** and **rt**
- $\text{PC} \leftarrow \text{PC} + 4$ calculate the address of the next instruction

All other R-type instructions will be similar.

The addi instruction

addi rs, rt, imm16

- $\text{mem}[\text{PC}]$ Fetch the instruction from memory
- $\text{R}[\text{rt}] \leftarrow \text{R}[\text{rs}] + \text{SignExt}(\text{imm16})$ Set register **rt** to the value of the sum of the contents of register **rs** and the immediate data word **imm16**
- $\text{PC} \leftarrow \text{PC} + 4$ calculate the address of the next instruction

All immediate arithmetic and logical instructions will be similar.

The **load** instruction

lw *rs*, *rt*, *imm16*

- $\text{mem}[\text{PC}]$
- $\text{Addr} \leftarrow \text{R}[\text{rs}] + \text{SignExt}(\text{imm16})$
- $\text{R}[\text{rt}] \leftarrow \text{Mem}[\text{Addr}]$
- $\text{PC} \leftarrow \text{PC} + 4$

Fetch the instruction from memory
Set memory address to the value of the sum of the contents of register **rs** and the immediate data word **imm16**

load the data at address **Addr** into register **rt**

calculate the address of the next instruction

The **store** instruction

sw *rs*, *rt*, *imm16*

- $\text{mem}[\text{PC}]$
- $\text{Addr} \leftarrow \text{R}[\text{rs}] + \text{SignExt}(\text{imm16})$
- $\text{Mem}[\text{Addr}] \leftarrow \text{R}[\text{rt}]$
- $\text{PC} \leftarrow \text{PC} + 4$

Fetch the instruction from memory
Set memory address to the value of the sum of the contents of register **rs** and the immediate data word **imm16**

store the data from register **rt** into memory at address **Addr**

calculate the address of the next instruction

The **branch** instruction

beq *rs*, *rt*, *imm16*

- $\text{mem}[\text{PC}]$ Fetch the instruction from memory
- $\text{Cond} \leftarrow \text{R}[\text{rs}] - \text{R}[\text{rt}]$ Evaluate the branch condition
- if ($\text{Cond} \text{ eq } 0$) calculate the address of the next instruction
 $\text{PC} \leftarrow \text{PC} + 4 + (\text{SignExt}(\text{imm16}) \times 4)$
- else $\text{PC} \leftarrow \text{PC} + 4$

The **jump** instruction

j *target* *target* is a memory address

- $\text{mem}[\text{PC}]$ Fetch the instruction from memory
- $\text{PC} \leftarrow \text{PC} + 4$ increment PC by 4
- $\text{PC}\langle 31:2 \rangle \leftarrow \text{PC}\langle 31:28 \rangle$ replace low order 28 bits with
 concat $\text{I}\langle 25:0 \rangle \ll 2$ the low order 26 bits from the instruction left shifted by 2


The Instruction Fetch Unit

Note that all instructions require that the **PC** be incremented.

We will design a datapath which performs this function — the **Instruction Fetch Unit**.

Its operation is described by the following:

- $\text{mem}[\text{PC}]$ Fetch the instruction from memory
- $\text{PC} \leftarrow \text{PC} + 4$ Increment the PC



Note that this does not yet handle branches or jumps.

Since it is the same for all instructions, when describing individual instructions this component will normally be omitted.

Datapath for R-type instructions

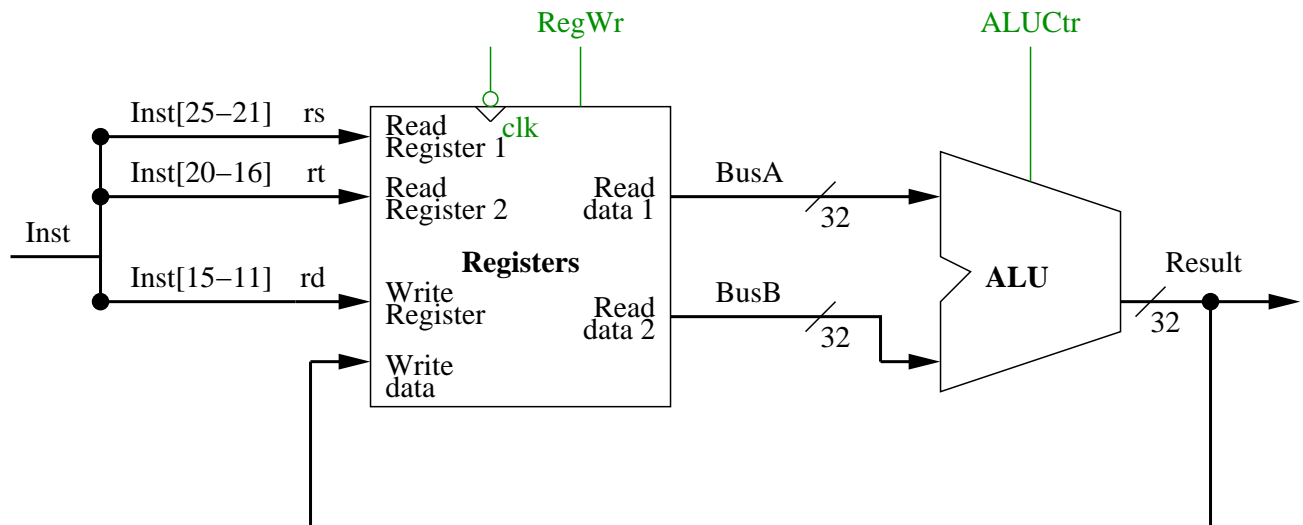
- $R[rd] \leftarrow R[rs] \text{ op } R[rt]$ Example: `add rd, rs, rt`

Recall that this instruction type has the following format:

R-type (register)

31	26	25	21	20	16	15	11	10	6	5	0				
op						rs		rt		rd		shamt		funct	
6 bits						5 bits		5 bits		5 bits		5 bits		6 bits	

The datapath contains the 32 bit register file and and ALU capable of performing all the required arithmetic and logic functions.



Note that the register is read from and written to at the “same time.” This implies that the register’s memory elements must be edge triggered, or are read and written on different clock phases, to allow the arithmetic operation to complete before the data is written in the register.

This datapath contains everything required to implement the required instructions **add**, **sub**, **and**, **or**, **slt**. All that is required is that the appropriate values be provided for the **ALUCtr** input for the required operation.

The register operands in the instruction field determine the registers which are read from and written to, and the **funct** field of the instruction determine which particular ALU operation is executed.

Recalling the control inputs for the ALU seen earlier, the values for the control input are:

ALU control lines	Function
000	and
001	or
010	add
110	subtract
111	set on less than

A control unit for the processor will be designed later.


It will set all the required control signals for each instruction, depending both on the particular instruction being executed (the **op code**) and, for r-type instructions, the **funct** field.

Datapath for Immediate arithmetic and logical instructions

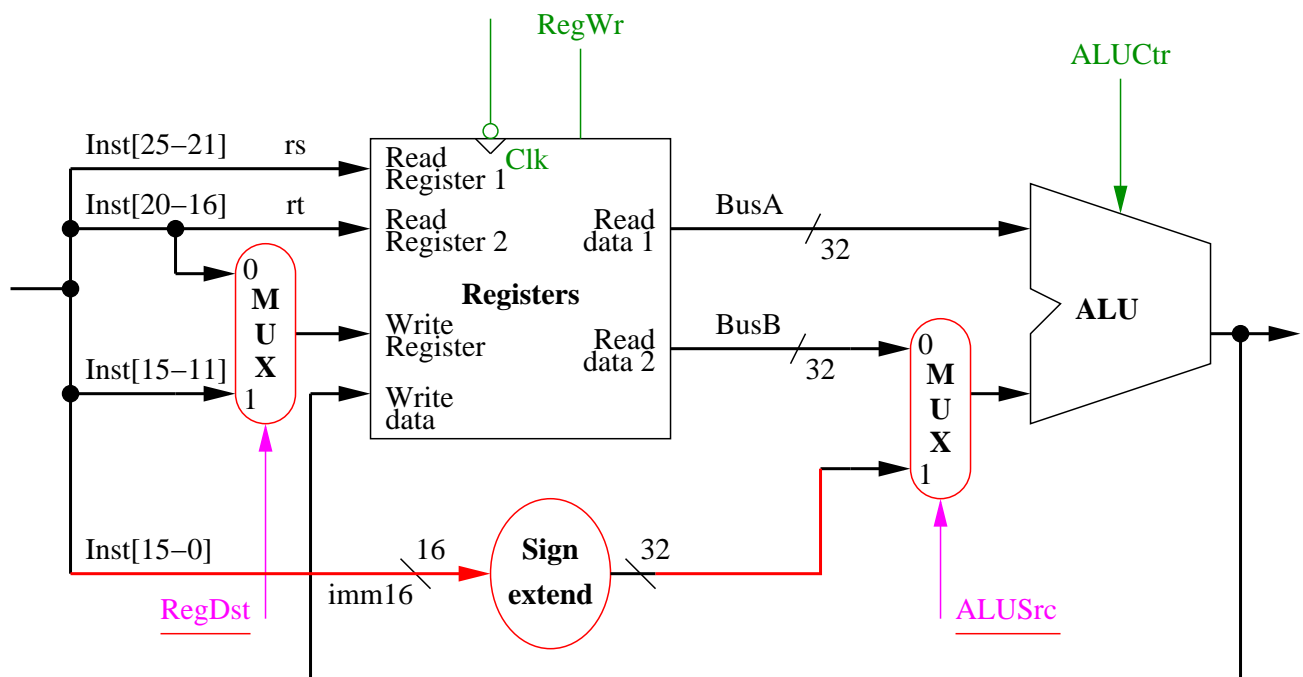
- $R[rt] \leftarrow R[rs] \text{ op } \text{imm16}$ Example: `addi rt, rs, imm16`

Recall that this instruction type has the following format:

I-type (immediate)



The main difference between this and an r-type instruction is that here one operand is taken from the instruction, and sign extended (for signed data) or zero extended (for logical and unsigned operations.)



Note the use of MUX's (with control inputs) to add functionality.

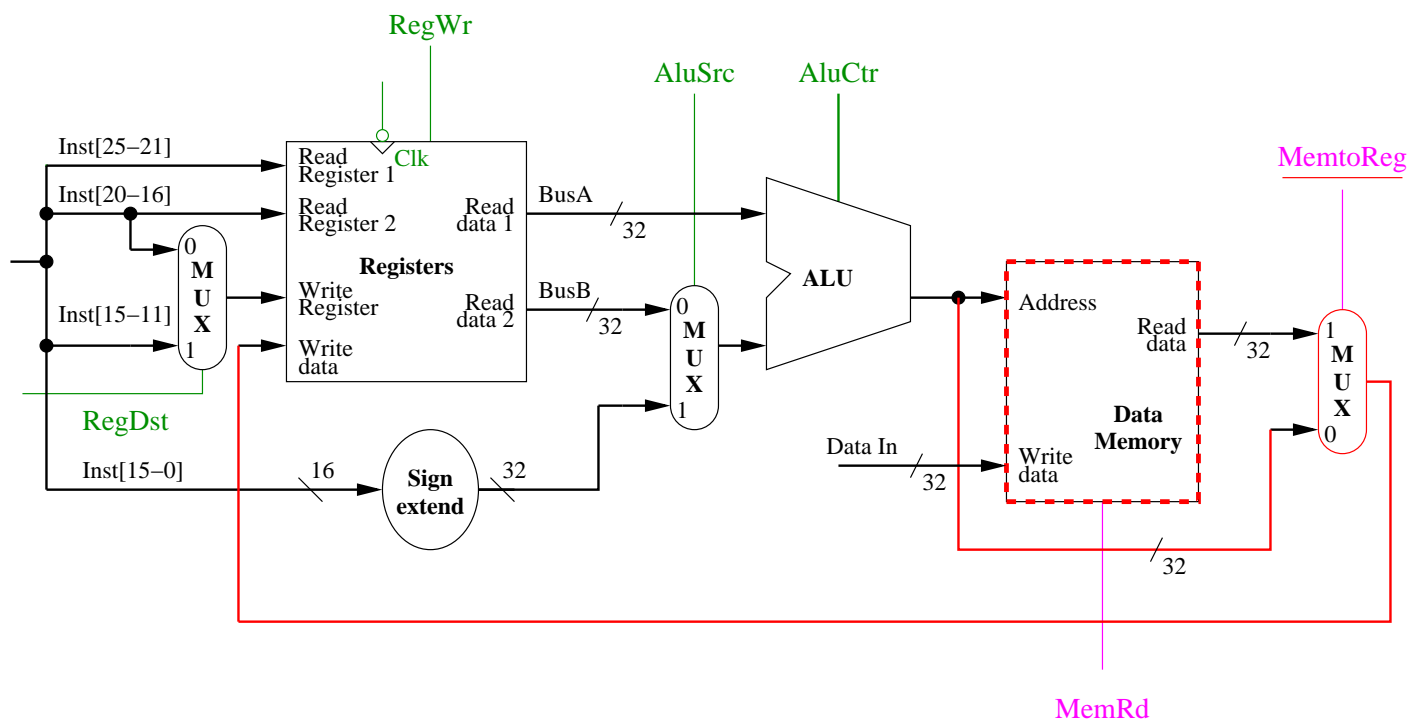

Datapath for the Load instruction

lw rt, rs, imm16

- $\text{Addr} \leftarrow \text{R}[\text{rs}] + \text{SignExt}(\text{imm16})$ Calculate the memory address
- $\text{R}[\text{rt}] \leftarrow \text{Mem}[\text{Addr}]$ load the data into register **rt**

This is also an immediate type instruction:

I-type (immediate)



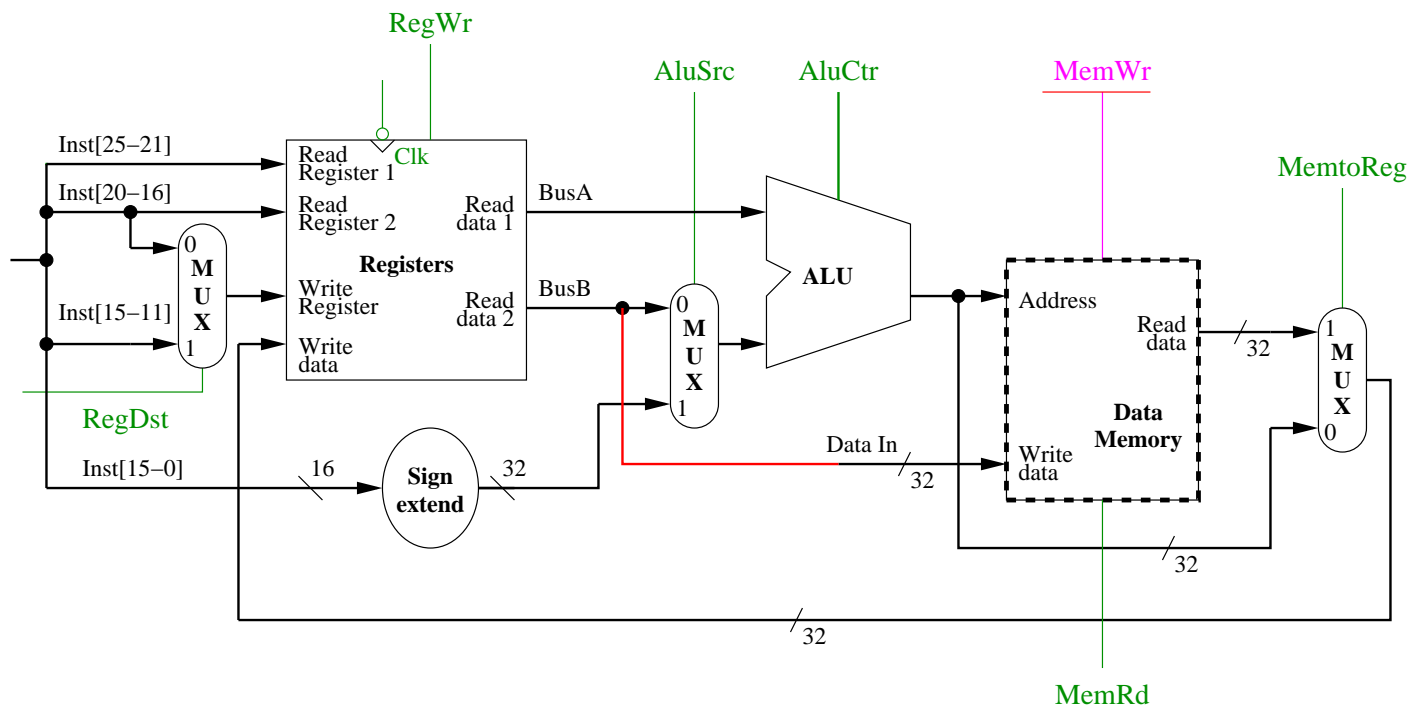

Datapath for the Store instruction

sw rt, rs, imm16

- $\text{Addr} \leftarrow R[\text{rs}] + \text{SignExt}(\text{imm16})$ Calculate the memory address
- $\text{Mem}[\text{Addr}] \leftarrow R[\text{rt}]$ Store the data from register **rt** to memory

This is also an immediate type instruction:

I-type (immediate)



Datapath for the Branch instruction

`beq rt, rs, imm16`

- $\text{Cond} \leftarrow R[\text{rs}] - R[\text{rt}]$ Calculate the branch condition
- if ($\text{Cond} \text{ eq } 0$) calculate the address of the next instruction
 $\text{PC} \leftarrow \text{PC} + 4 + (\text{SignExt}(\text{imm16}) \times 4)$
- else $\text{PC} \leftarrow \text{PC} + 4$

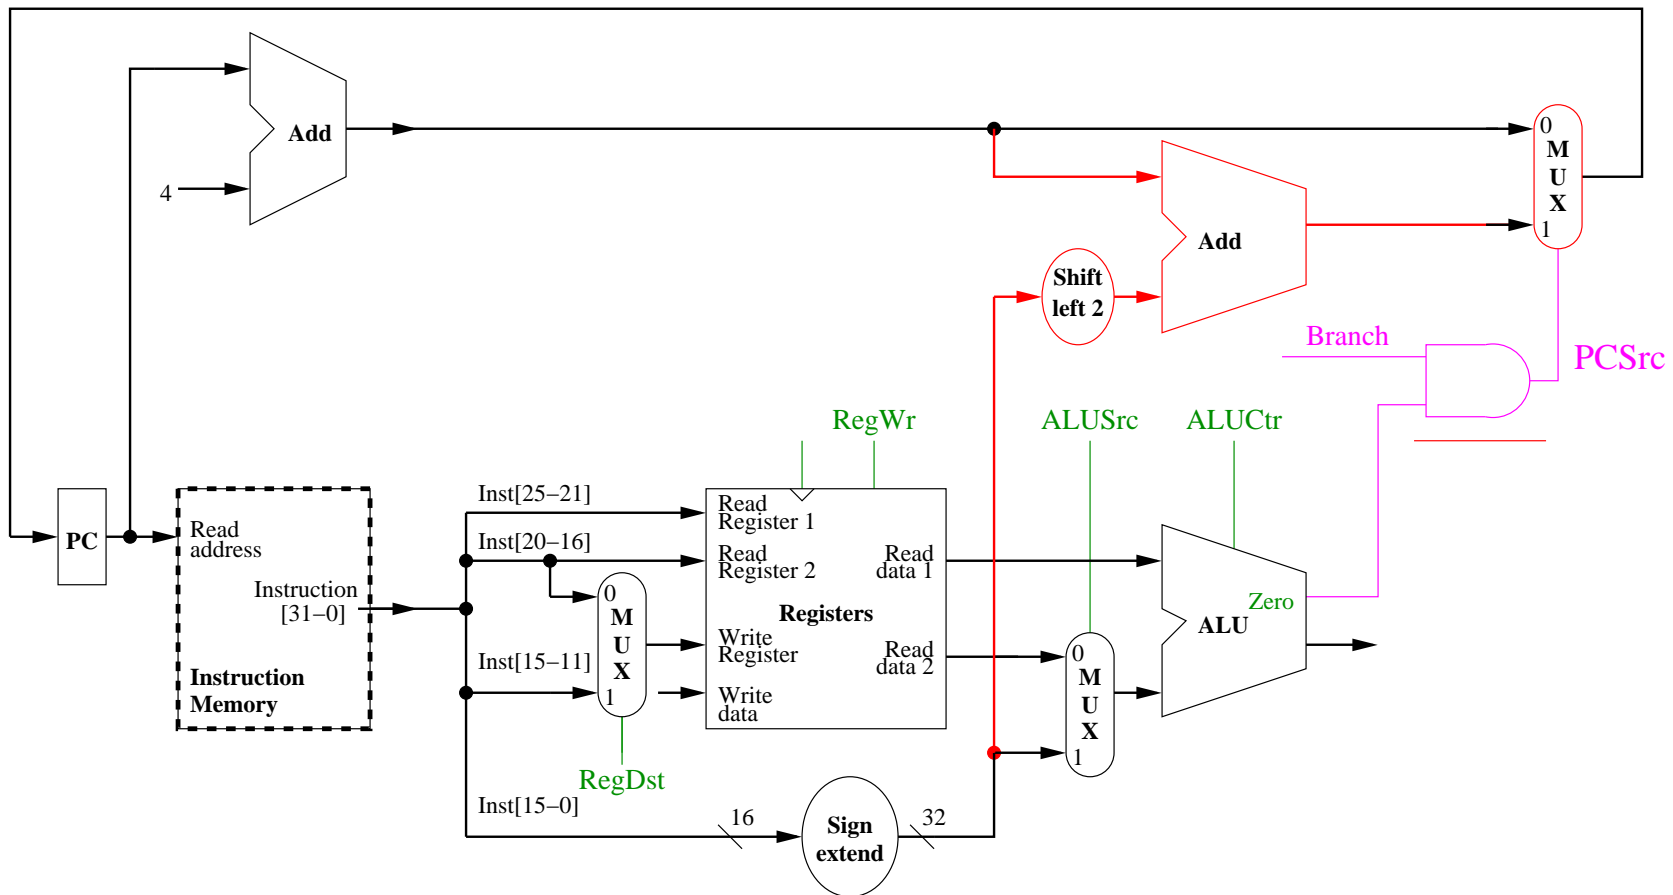
This is also an immediate type instruction.

In the **load** and **store** instructions, the ALU was used to calculate the address for data memory.

It is possible to do this for the **branch** instructions as well, but it would require first performing the comparison using the ALU, and then using the ALU to calculate the address.

This would require two clock periods, in order to sequence the operations correctly.

A faster implementation would be to provide another adder to implement the address calculation. This is what we will do, for the present example.



Datapath for the Jump instruction


j target

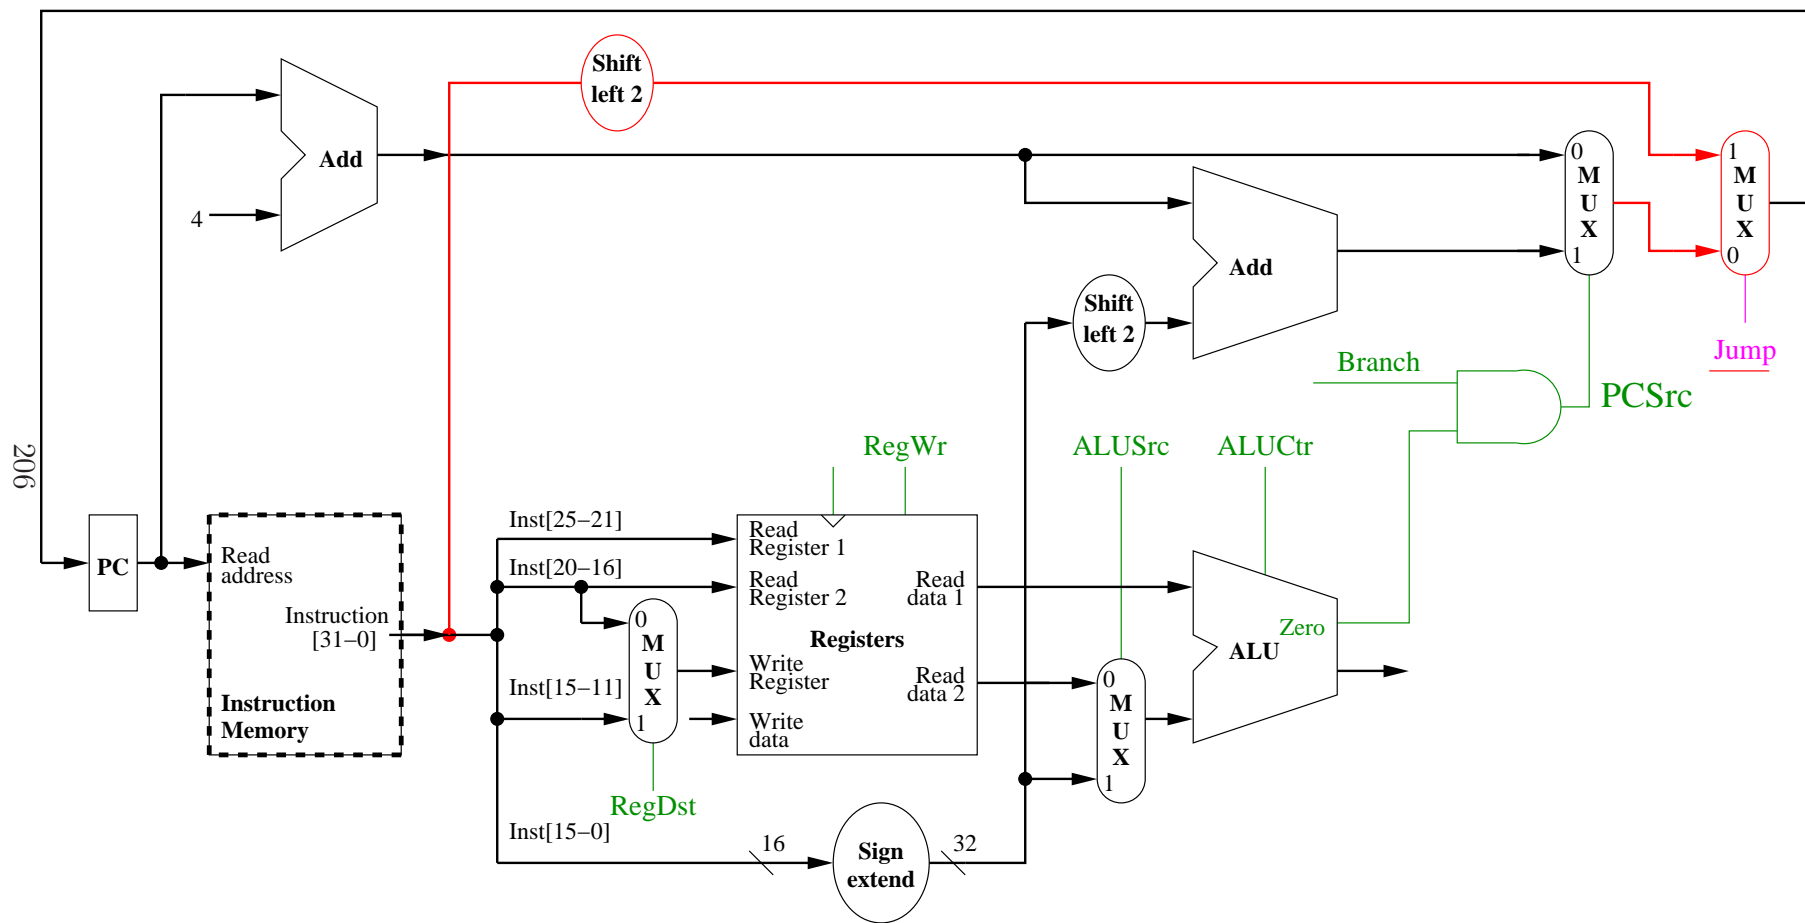
- $PC\langle 31:2 \rangle \leftarrow PC\langle 31:28 \rangle \text{ concat target}\langle 25:0 \rangle$ Calculate the jump address by concatenating the high order 4 bits of the PC with the target address

Here, the address calculation is just obtained from the high order 4 bits of the PC and the 26 bits (shifted left by 2 bits to make 28) of the target address.

The additions to the datapath are straightforward.

J-type (jump)





Putting it together

The datapath was shown in segments, some of which built on each other.

Required control signals were identified, and all that remains is to:

1. Combine the datapath elements
2. Design the appropriate control signals

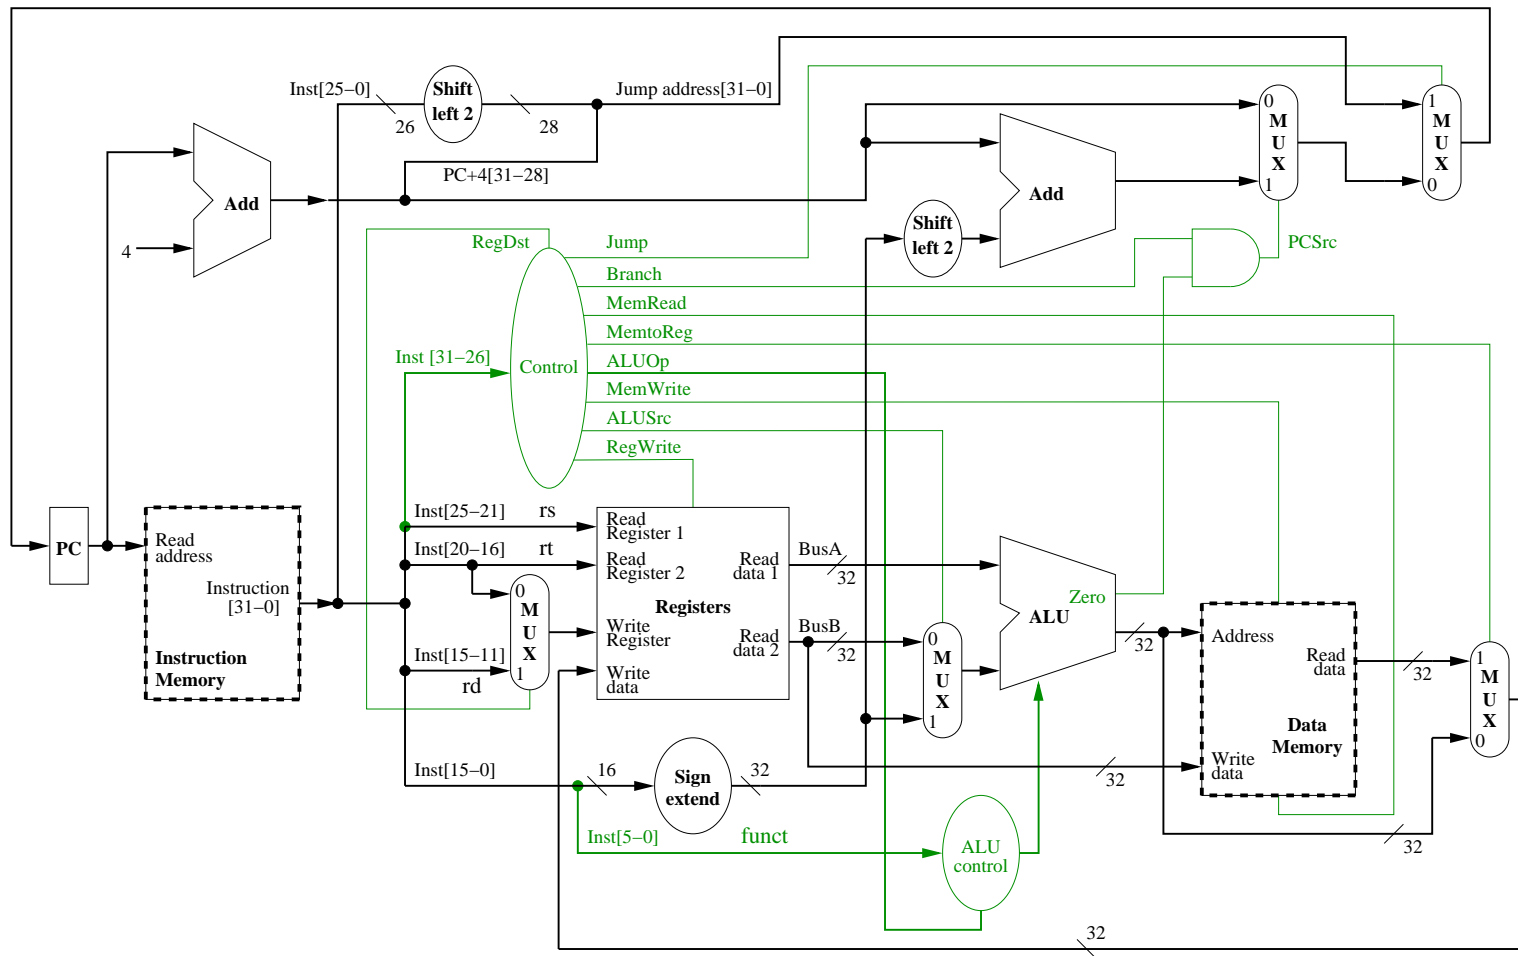
Combining the datapath elements is rather straightforward, since we have mainly built up the datapath by adding functionality to accommodate the different instruction types.

When two paths are required, we have implemented both and used multiplexors to choose the appropriate results.

The required control signals are mainly the inputs for those MUX's and the signals required by the ALU.

The next slide shows the combined data path, and the required control signals.

The actual control logic is yet to be designed.



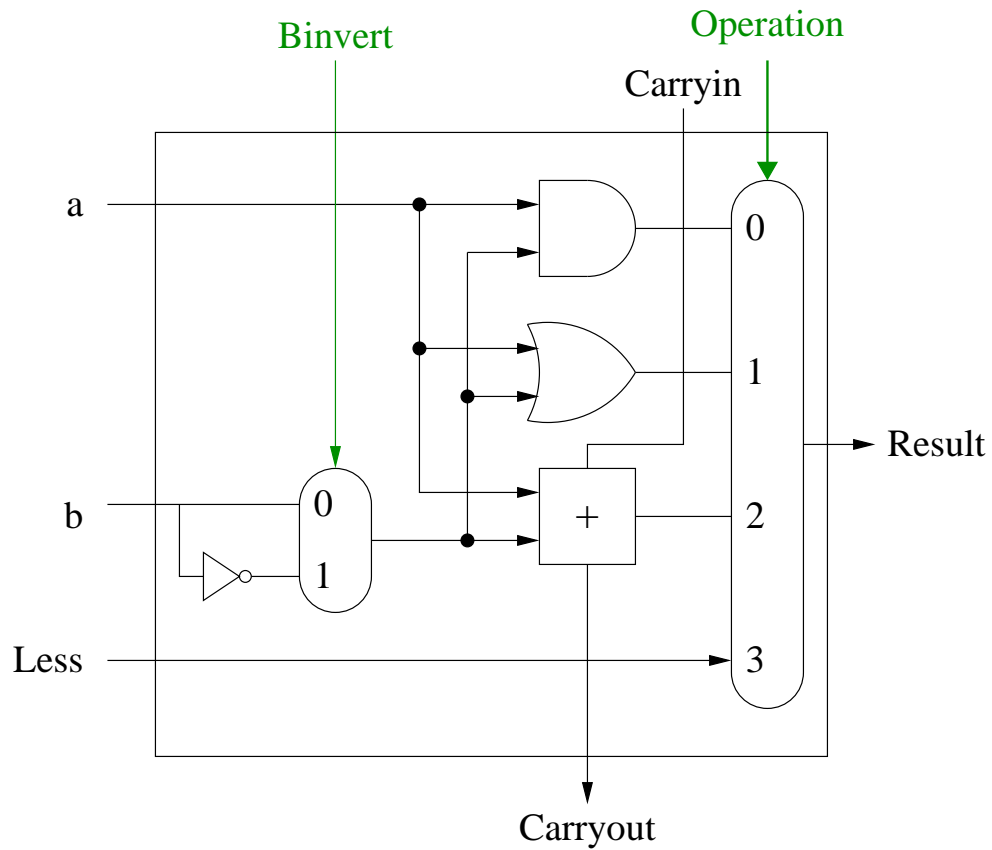
Designing the control logic

The control logic depends on the details of the devices in the control path, and on the individual bits in the op code for the instructions. The arithmetic and logic operations for the r-type instructions also depend on the **funct** field of the instruction.

The datapath elements we have used are:

- a 32 bit ALU with an output indicating if the result is zero
- adders
- MUX's (2 line to 1-line)
- a 32 register \times 32 bits/register register file
- individual 32 bit registers
- a sign extender
- instruction memory
- data memory

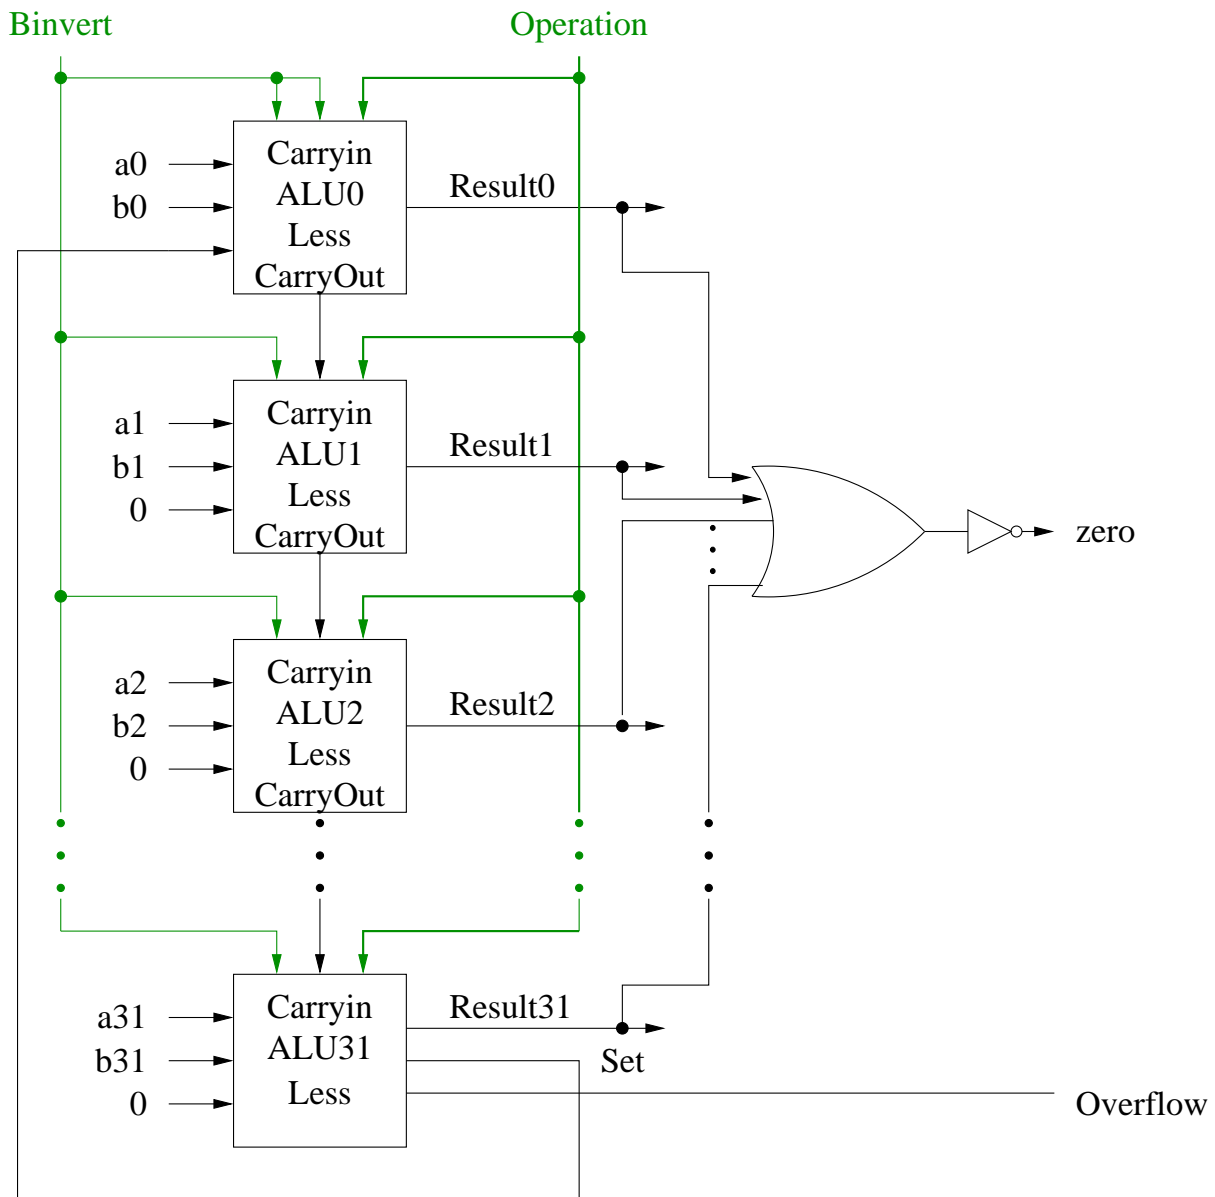
The ALU — a single bit



Note that there are three control bits; the single bit **Binvert**, and the two bit input to the MUX, labeled **Operation**.

The ALU performs the operations **and**, **or**, **add**, and **subtract**.

The 32 bit ALU



ALU control lines	Function
000	and
001	or
010	add
110	subtract
111	set on less than

We will design the control logic to implement the following instructions (others can be added similarly):

Name	Op-code					
	Op5	Op4	Op3	Op2	Op1	Op0
R-format	0	0	0	0	0	0
lw	1	0	0	0	1	1
sw	1	0	1	0	1	1
beq	0	0	0	1	0	0
j	0	0	0	0	1	0

Note that we have omitted the immediate arithmetic and logic functions.

The **funct** field will also have to be decoded to produce the required control signals for the ALU.

A separate decoder will be used for the main control signals and the ALU control. This approach is sometimes called *local decoding*. Its main advantage is in reducing the size of the main controller.

The control signals

The signals required to control the datapath are the following:

- **Jump** — set to 1 for a **jump** instruction
- **Branch** — set to 1 for a **branch** instruction
- **MemtoReg** — set to 1 for a **load** instruction
- **ALUSrc** — set to 0 for r-type instructions, and 1 for instructions using immediate data in the ALU (**beq** requires this set to 0)
- **RegDst** — set to 1 for r-type instructions, and 0 for immediate instructions
- **MemRead** — set to 1 for a **load** instruction
- **MemWrite** — set to 1 for a **store** instruction
- **RegWrite** — set to 1 for any instruction writing to a register
- **ALUOp** (k bits) — encodes ALU operations except for r-type operations, which are encoded by the **funct** field

For the instructions we are implementing, **ALUOp** can be encoded using 2 bits as follows:

ALUOp[1]	ALUOp[0]	Instruction
0	0	memory operations (load , store)
0	1	beq
1	0	r-type operations

The following tables show the required values for the control signals as a function of the instruction op codes:

Instruction	Op-code	RegDst	ALUSrc	MemtoReg	Reg Write
r-type	0 0 0 0 0 0	1	0	0	1
lw	1 0 0 0 1 1	0	1	1	1
sw	1 0 1 0 1 1	x	1	x	0
beq	0 0 0 1 0 0	x	0	x	0
j	0 0 0 0 1 0	x	x	x	0

Instruction	Op-code	Mem Read	Mem Write	Branch	ALUOp[1:0]	Jump
r-type	0 0 0 0 0 0	0	0	0	1 0	0
lw	1 0 0 0 1 1	1	0	0	0 0	0
sw	1 0 1 0 1 1	0	1	0	0 0	0
beq	0 0 0 1 0 0	0	0	1	0 1	0
j	0 0 0 0 1 0	0	0	0	x x	1

This is all that is required to implement the control signals; each control signal can be expressed as a function of the op-code bits.

For example,

$$\text{RegDst} = \overline{\text{Op5}} \cdot \overline{\text{Op4}} \cdot \overline{\text{Op3}} \cdot \overline{\text{Op2}} \cdot \overline{\text{Op1}} \cdot \overline{\text{Op0}}$$

$$\text{ALUSrc} = \text{Op5} \cdot \overline{\text{Op4}} \cdot \overline{\text{Op2}} \cdot \text{Op1} \cdot \text{Op0}$$

All that remains is to design the control for the ALU.

The ALU control

The inputs to the ALU control are the **ALUOp** control signals, and the 6 bit **funct** field.

The **funct** field determines the ALU operations for the r-type operations, and **ALUOp** signals determine the ALU operations for the other types of instructions.

Previously, we saw that if **ALUOp**[1] was 1, it indicated an r-type operation. **ALUOp**[0] was set to 0 for memory operations (requiring the ALU to perform an **add** operation to calculate the address for data) and to 1 for the **beq** operation, requiring a subtraction to compare the two operands.

The ALU itself requires three inputs.

The following table shows the required inputs and outputs for the instructions using the ALU:

Instruction	ALUOp	funct	ALU operation	ALU control input
lw	0 0	x x x x x x	add	0 1 0
sw	0 0	x x x x x x	add	0 1 0
beq	0 1	x x x x x x	subtract	1 1 0
add	1 0	1 0 0 0 0 0	add	0 1 0
sub	1 0	1 0 0 0 1 0	subtract	1 1 0
and	1 0	1 0 0 1 0 0	AND	0 0 0
or	1 0	1 0 0 1 0 1	OR	0 0 1
slt	1 0	1 0 1 0 1 0	set on less than	1 1 1

The time required for single cycle instructions

Arithmetic and logical instructions

PC	Inst. Memory	Reg. Read	mux	ALU	mux	Reg. Write
----	--------------	-----------	-----	-----	-----	------------

time →

Branch

PC	Inst. Memory	Reg. Read	mux	ALU	mux	mux
			Sign ext.	add		

(The sign extension and add occur in parallel with the other operations, register read and ALU comparison)

Load

PC	Inst. Memory	Reg. Read	mux	ALU	Data Memory	mux	Reg. Write
----	--------------	-----------	-----	-----	-------------	-----	------------



The "critical path"

Store

PC	Inst. Memory	Reg. Read	mux	ALU	Data Mem.
----	--------------	-----------	-----	-----	-----------


Jump

PC	Inst. Memory	mux
----	--------------	-----

The clock period must be at least as long as the time for the critical path.


Looking back at the instruction timing for the single cycle processor, we see that the **load** instruction requires two memory accesses, and therefore will require at least two cycles.

Arithmetic and logical instructions



time →

Load




The "critical path"

Jump



Considering the option of using the ALU to increment the PC, note also that if the PC is read at the beginning of a cycle and loaded at the end of the cycle, then it can be incremented in parallel with the memory access. Also, if the diagram really represents the time for the various operations, the register and MUX operations together require approximately the same time as a memory operation, requiring *five* cycles in total.



The "critical path"

A multi-cycle implementation

We will consider the design of a multi-cycle implementation of the processor developed so far. The processor will have:

- a single memory for instructions and data
- a single ALU for both addressing and data operations
- instructions requiring different numbers of cycles

There are now resource limitations — only one access to memory, one access to the register file, and one ALU operation can occur in each clock cycle.

It is clear that both the instruction and data would be required during the execution of an instruction. Additional registers, the **instruction register (IR)** and the **memory data register (MDR)** will be required to hold the instruction and data words from memory between cycles.


Registers may also be required to hold the register operands from **BusA** and **BusB** (registers **A** and **B**, respectively).

(Recall that the **branch** instructions require an arithmetic comparison *before* an address calculation.)

We will look at each type of instruction individually to determine if it can actually be done with the time and resources available.

The R-type instructions

- $R[rd] \leftarrow R[rs] \text{ op } R[rt]$ Example: `add rd, rs, rt`



Clearly, the instruction can be completed in four cycles, from the timing. We need only determine if the required resources are available.

- In the first cycle, the instruction is fetched from memory, and the ALU is used to increment the PC. The instruction must be saved in the instruction register (**IR**) so it can be used in the following cycles. (This may extend the cycle time).
- In the second cycle, the registers are read, and the values from the registers to be used by the ALU must be saved, in registers **A** and **B**, again new registers.
- In the third cycle, the r-type operation is completed in the ALU, and the result saved in another new register, **ALUOut**.
- In the fourth cycle, the value in register **ALUOut** is written into the register file.

Four registers had to be added to preserve values from one cycle to the next, but there were no resource conflicts — the ALU was required only in the first and third cycle.

We can capture these steps in an RTL description:

Cycle 1	$IR \leftarrow mem[PC]$	Save instruction in IR
	$PC \leftarrow PC + 4$	increment PC
Cycle 2	$A \leftarrow R[rs]$	save register values for next cycle
	$B \leftarrow R[rt]$	
Cycle 3	$ALUOut \leftarrow A \text{ op } B$	calculate result and store in ALUOut
Cycle 4	$R[rd] \leftarrow ALUOut$	store result in register file


This is really an expansion of the original RTL description of the R-type instructions, where the internal registers are also used. The original description was:

$mem[PC]$	Fetch the instruction from memory
$R[rd] \leftarrow R[rs] \text{ op } R[rt]$	Set register rd to the value of the operation applied to the contents of registers rs and rt
$PC \leftarrow PC + 4$	calculate the address of the next instruction

When using a “silicon compiler” to design a processor, designers often refine the RTL description in a similar way in order to achieve a more efficient implementation for the datapath or control.

The Branch instruction — beq

- $\text{Cond} \leftarrow R[\text{rs}] - R[\text{rt}]$ Calculate the branch condition
- if ($\text{Cond} \text{ eq } 0$) calculate the address of the next instruction
 - $\text{PC} \leftarrow \text{PC} + 4 + (\text{SignExt}(\text{imm16}) \times 4)$
- else $\text{PC} \leftarrow \text{PC} + 4$



In this case, *three* arithmetic operations are required, (incrementing the PC, comparing the register values, and adding the immediate field to the PC.)

Clearly, the comparison could not be done until the values have been read from the register, so this must be done in cycle 3.

The address calculation could be done in cycle 2, however, since it uses only data from the instruction (the immediate field) and the new value of the PC, and the ALU is not being used in this cycle. The result would have to be stored in a register, to be used in the next cycle. We could use the register `ALUOut` for this, since the R-type operations only require it at the end of cycle 3.

Recall that the ALU produced an output **Zero** which could be used to implement the comparison. It is available during the third cycle, and could be used to enable the replacement of the PC with the value stored in `ALUOut` in the previous cycle.

The original RTL for the **beq** was:

- $\text{mem}[\text{PC}]$ Fetch the instruction from memory
- $\text{Cond} \leftarrow \text{R}[\text{rs}] - \text{R}[\text{rt}]$ Evaluate the branch condition
- if ($\text{Cond} \text{ eq } 0$) calculate the address of the next instruction
 $\text{PC} \leftarrow \text{PC} + 4 + (\text{SignExt}(\text{imm16}) \times 4)$
- else $\text{PC} \leftarrow \text{PC} + 4$

Rewriting the RTL code for the **beq** instruction, including the operations on the internal registers, we have:


- Cycle 1 $\text{IR} \leftarrow \text{mem}[\text{PC}]$ Save instruction in **IR**
 $\text{PC} \leftarrow \text{PC} + 4$ increment **PC**
- Cycle 2 $\text{A} \leftarrow \text{R}[\text{rs}]$ save register values for next cycle
 $\text{B} \leftarrow \text{R}[\text{rt}]$ (for comparison)
 $\text{ALUOut} \leftarrow \text{PC} +$ calculate address for branch
 $\text{signextend}(\text{imm16}) \ll 2$ and place in **ALUOut**
- Cycle 3 Compare **A** and **B**
 if **Zero** is set replace **PC** with **ALUOut** if **Zero**
 then $\text{PC} \leftarrow \text{ALUOut}$ is set, otherwise do not change **PC**

Note that this instruction now requires three cycles.

Also, the first cycle is identical to that of the R-type instructions. The second cycle does the same as the R-type, and also does the address calculation. Note that, at this point, the instruction may not require the result of the address calculation, but it is calculated anyway.

The Load instruction

- $\text{Addr} \leftarrow \text{R}[\text{rs}] + \text{SignExt}(\text{imm16})$ Calculate the memory address
- $\text{R}[\text{rt}] \leftarrow \text{Mem}[\text{Addr}]$ load data into register **rt**



Clearly, the first cycle is the same as in the previous examples.

For the second cycle, register **R[rs]** contains part of an address, and register **R[rt]** contains a value to be saved in memory (for **store**) or to be replaced from memory (for **load**). They must therefore be saved in registers (**A** and **B**) for future use, like the previous instructions.

In the third cycle, the address is calculated from the contents of **A** and the **imm16** field of the instruction and stored in a register (**ALUOut**) for use in the next cycle.

This address (now in **ALUOut**) is used to access the appropriate memory location in the fourth cycle, and the contents of memory are placed in a register **MDR**, the memory data register.

In the fifth cycle, the contents of the **MDR** are stored in the register file in register **R[rt]**.

The original RTL for `load` was:

- `mem[PC]` Fetch the instruction from memory
- `Addr \leftarrow R[rs] + SignExt(imm16)` Set memory address to the value of the sum of the contents of register `rs` and the immediate data word `imm16`
- `R[rt] \leftarrow Mem[Addr]` load the data at address `Addr` into register `rt`
- `PC \leftarrow PC + 4` calculate the address of the next instruction


The RTL for this implementation is:

- | | | |
|---------|---|---|
| Cycle 1 | <code>IR \leftarrow mem[PC]</code> | Save instruction in <code>IR</code> |
| | <code>PC \leftarrow PC + 4</code> | increment <code>PC</code> |
| Cycle 2 | <code>A \leftarrow R[rs]</code> | save address register for next cycle |
| | <code>B \leftarrow R[rt]</code> | |
| Cycle 3 | <code>ALUOut \leftarrow A + signextend(imm16)</code> | calculate address for data and place in <code>ALUOut</code> |
| Cycle 4 | <code>MDR \leftarrow Mem[ALUOut]</code> | store contents of memory at address <code>ALUOut</code> in <code>MDR</code> |
| Cycle 5 | <code>R[rt] \leftarrow MDR</code> | store value originally from memory in <code>R[rt]</code> |

Recall that this instruction was the longest instruction in the single cycle implementation.

The Store instruction

- $\text{Addr} \leftarrow \text{R}[\text{rs}] + \text{SignExt}(\text{imm16})$ Calculate the memory address
- $\text{Mem}[\text{Addr}] \leftarrow \text{R}[\text{rt}]$ store the contents of register **rt** in memory




The **store** instruction is much like the **load** instruction, except that the value in register **R[rt]** is written into memory, rather than read from it.

The main difference is that, in the fourth cycle, the address calculated from **R[rs]** and **imm16** (and saved in **ALUOut**) is used to store the value from register **R[rt]** in memory.

A fifth cycle is not required.

The Jump instruction

- $PC\langle 31:2 \rangle \leftarrow PC\langle 31:28 \rangle$ Calculate the jump address by concatenating the high order 4 bits of the PC with the target address
 $\text{concat target}\langle 25:0 \rangle$



The first cycle, which fetches the instruction from memory and places it in **IR**, and increments **PC** by 4, is the same as other instructions.

The next operation is to concatenate the low order 26 bits of the instruction with the high order 4 bits of the **PC**.

In the PC, the low order 2 bits are 0, so they are not actually loaded or stored.

The shift of the bits from the instruction can be accomplished without any additional hardware, merely by connecting bit **IR[25]** to bit **PC[27]**, etc.

Note that adding 4 to the **PC** may cause the four high order bits to change.

Could this cause problems ?

Changes to the datapath for a multi-cycle implementation

We have found that several additional registers are required in the multi-cycle datapath in order to save information from one cycle to the next.

These were the registers **IR**, **MDR**, **A**, **B**, and **ALUOut**.

The overall hardware complexity may be reduced, however, since the adders required for addressing have been replaced by the ALU.

Recall that the primary reason for choosing five cycles was the assumption that the time to obtain a value from memory was the single slowest operation in the datapath. Also, we assumed that the register file operations take a smaller, but comparable, amount of time.

If either of these conditions were not true, then quite a different schedule of operations might have been chosen.

The datapath for the multi-cycle processor

Fortunately, after our design of the single cycle processor, we have a good idea of the datapath elements required to implement each individual instruction. We can also seek opportunities to reuse functional blocks in different cycles, potentially reducing the number of hardware blocks (and hence the complexity and cost) of the datapath.

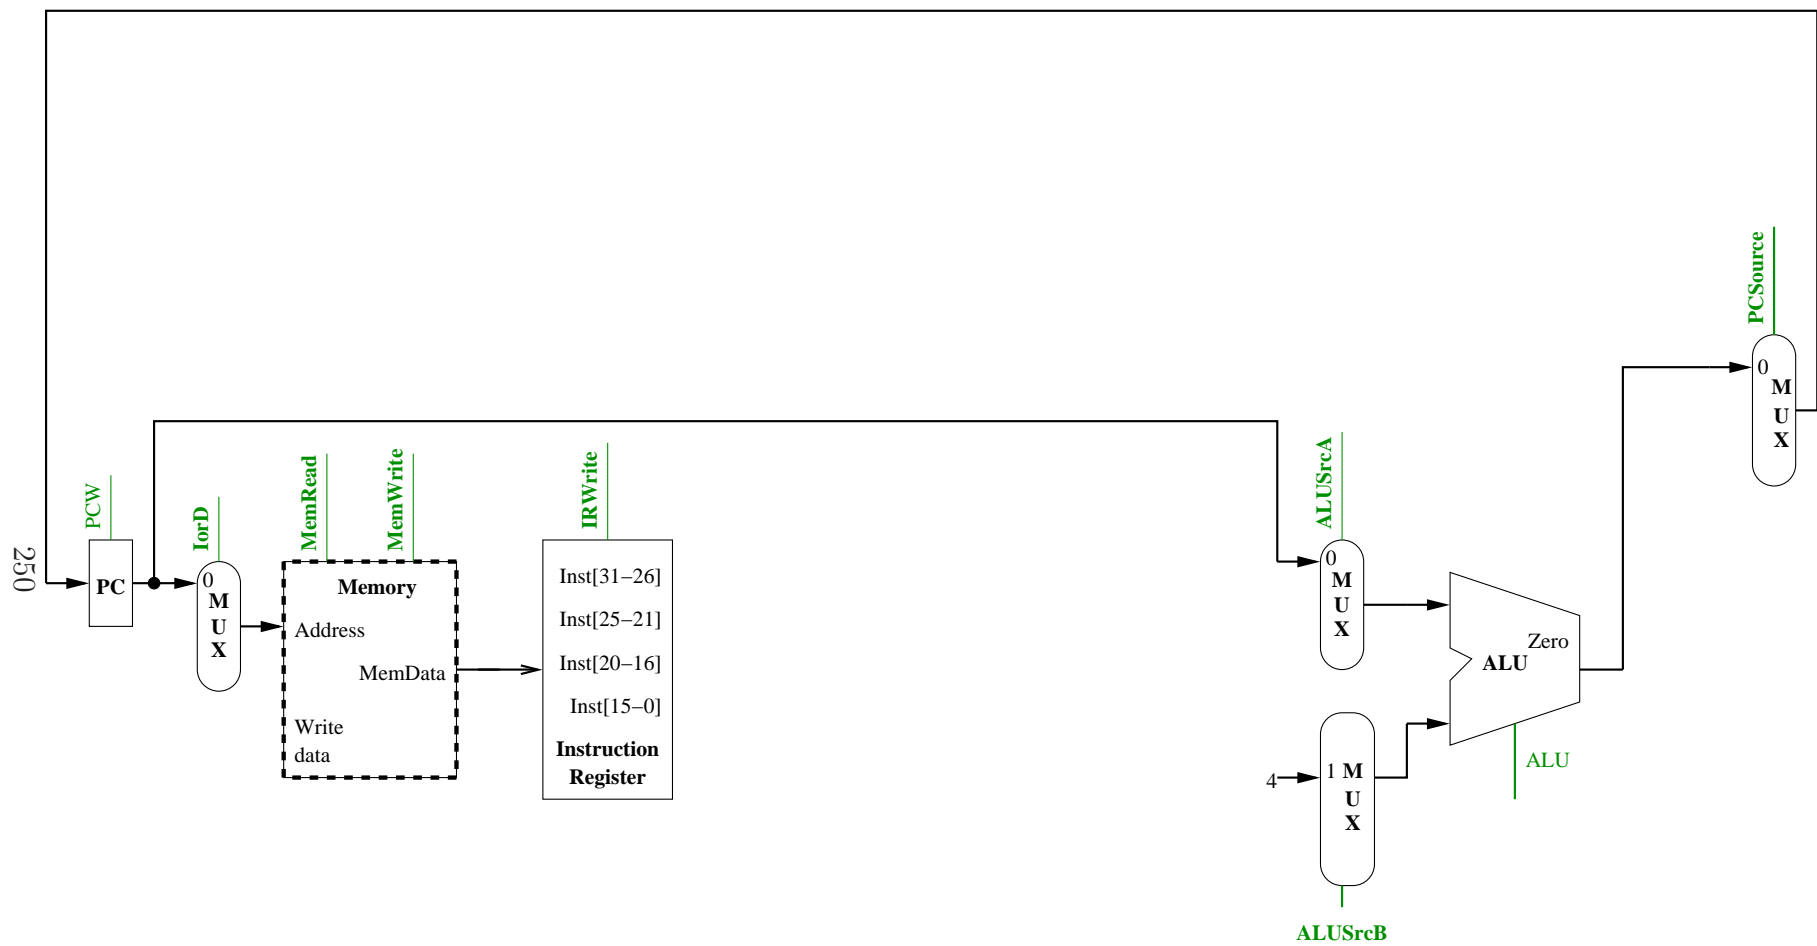
The datapath for the multi-cycle processor is similar to that of the single cycle processor, with

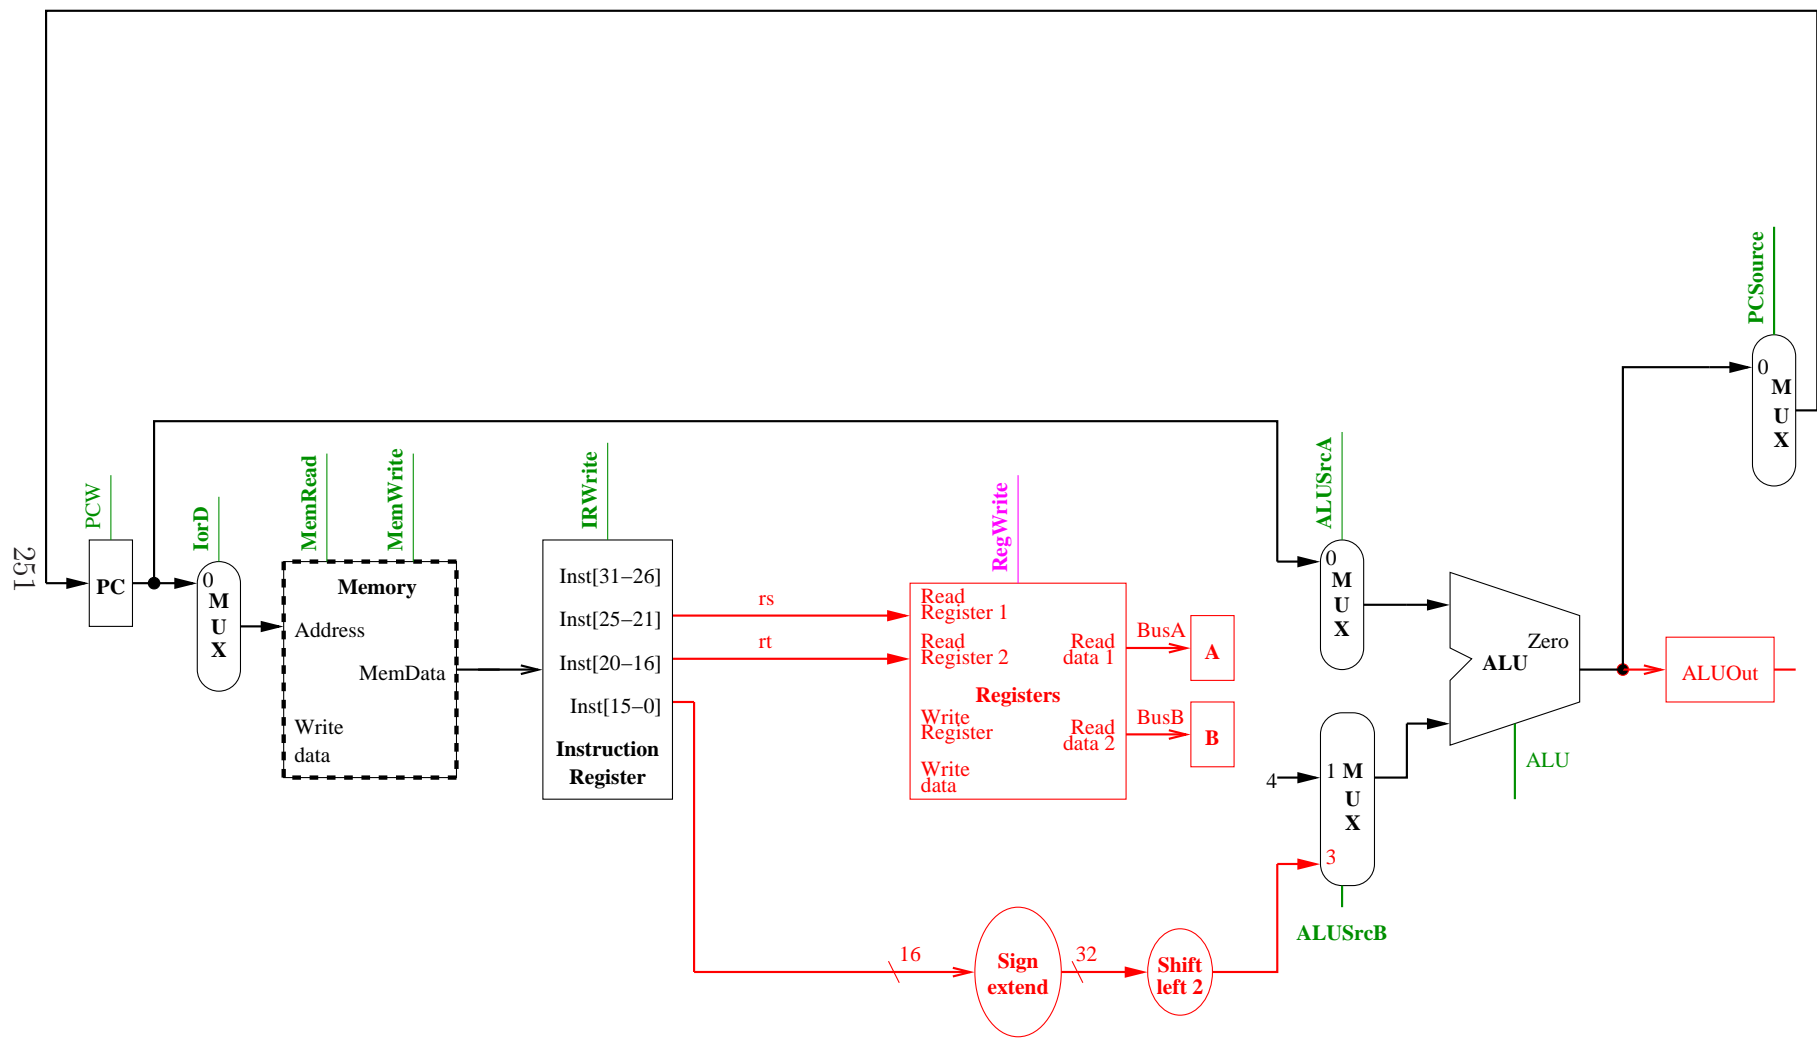
- the addition of the registers noted (**IR**, **MDR**, **A**, **B**, and **ALUOut**)
- the elimination of the adders for address calculation
- a MUX must be extended because there are now three separate calculations for the next address (**jump**, **branch**, and the normal incrementing of the PC).
- additional control signals controlling the writing of the registers.

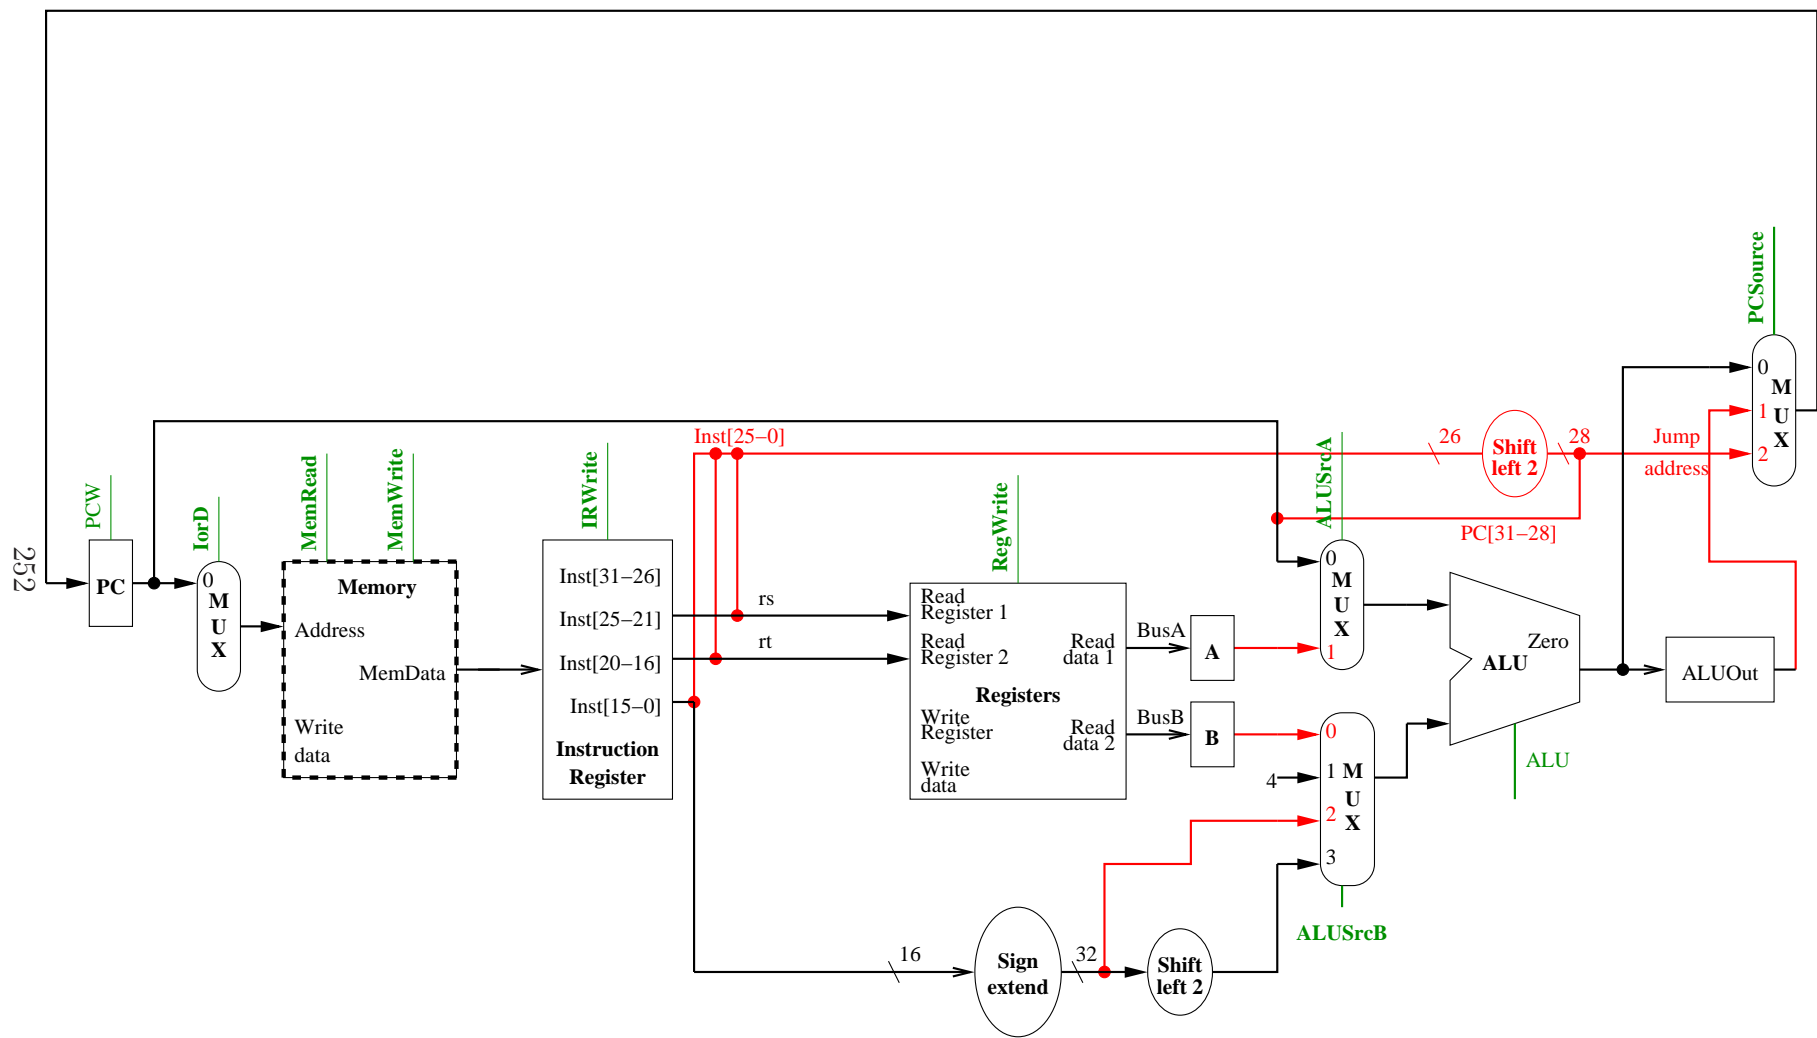
The following diagrams show the datapath for the multi-cycle implementation of the processor.

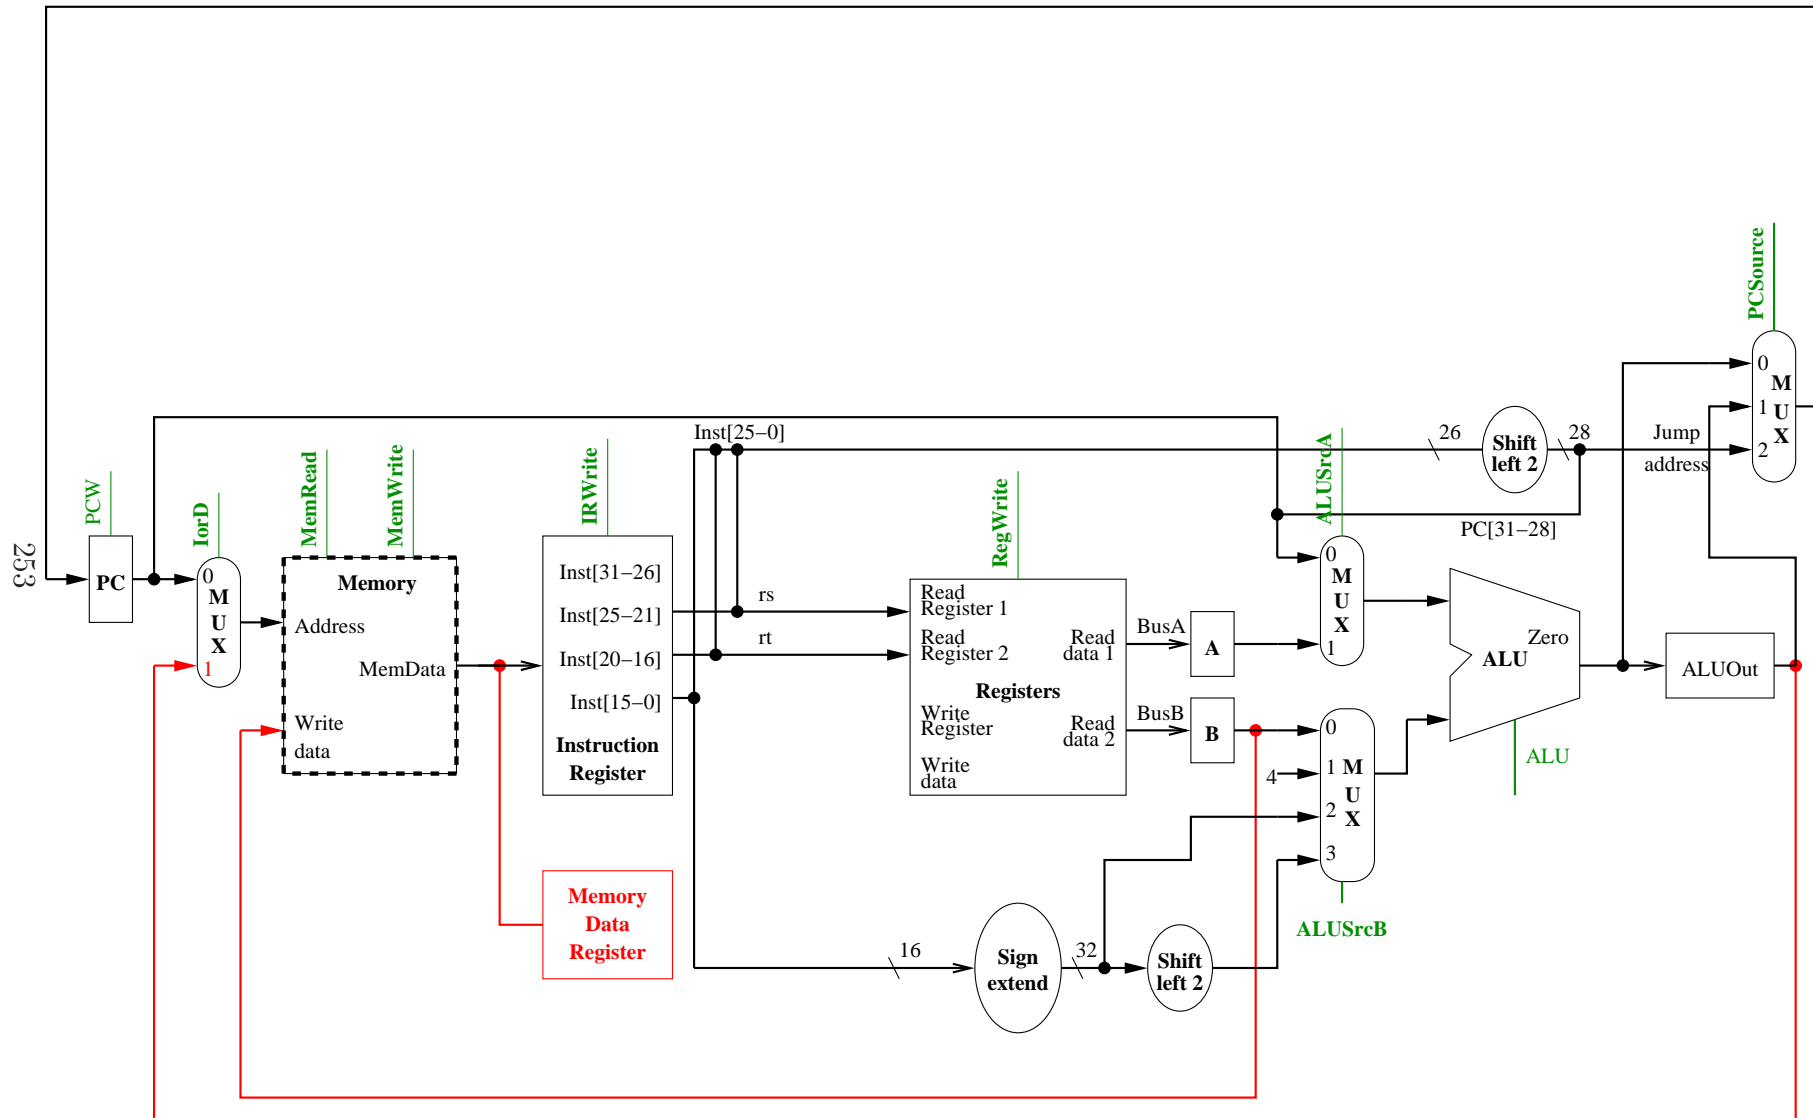
The additions to the datapath for each cycle is shown in red.

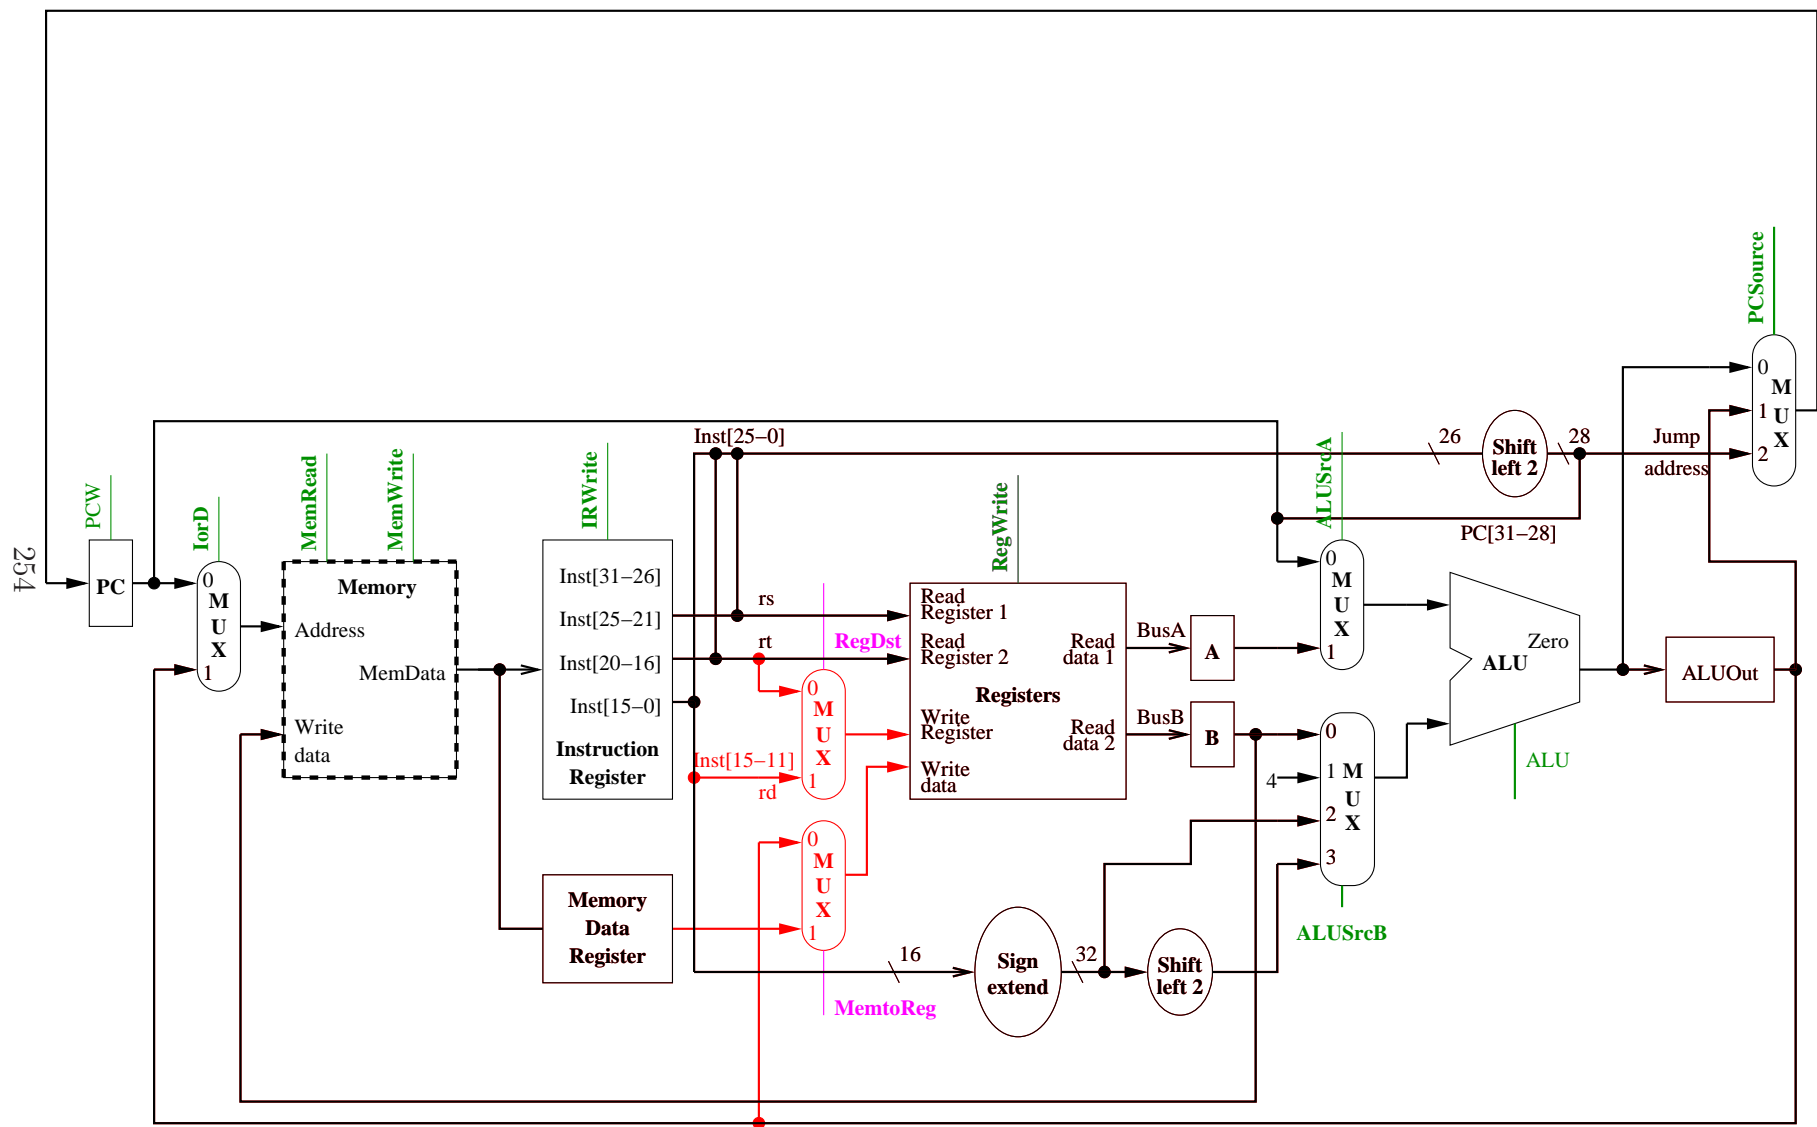
The required control signals are shown in green in the final figure.

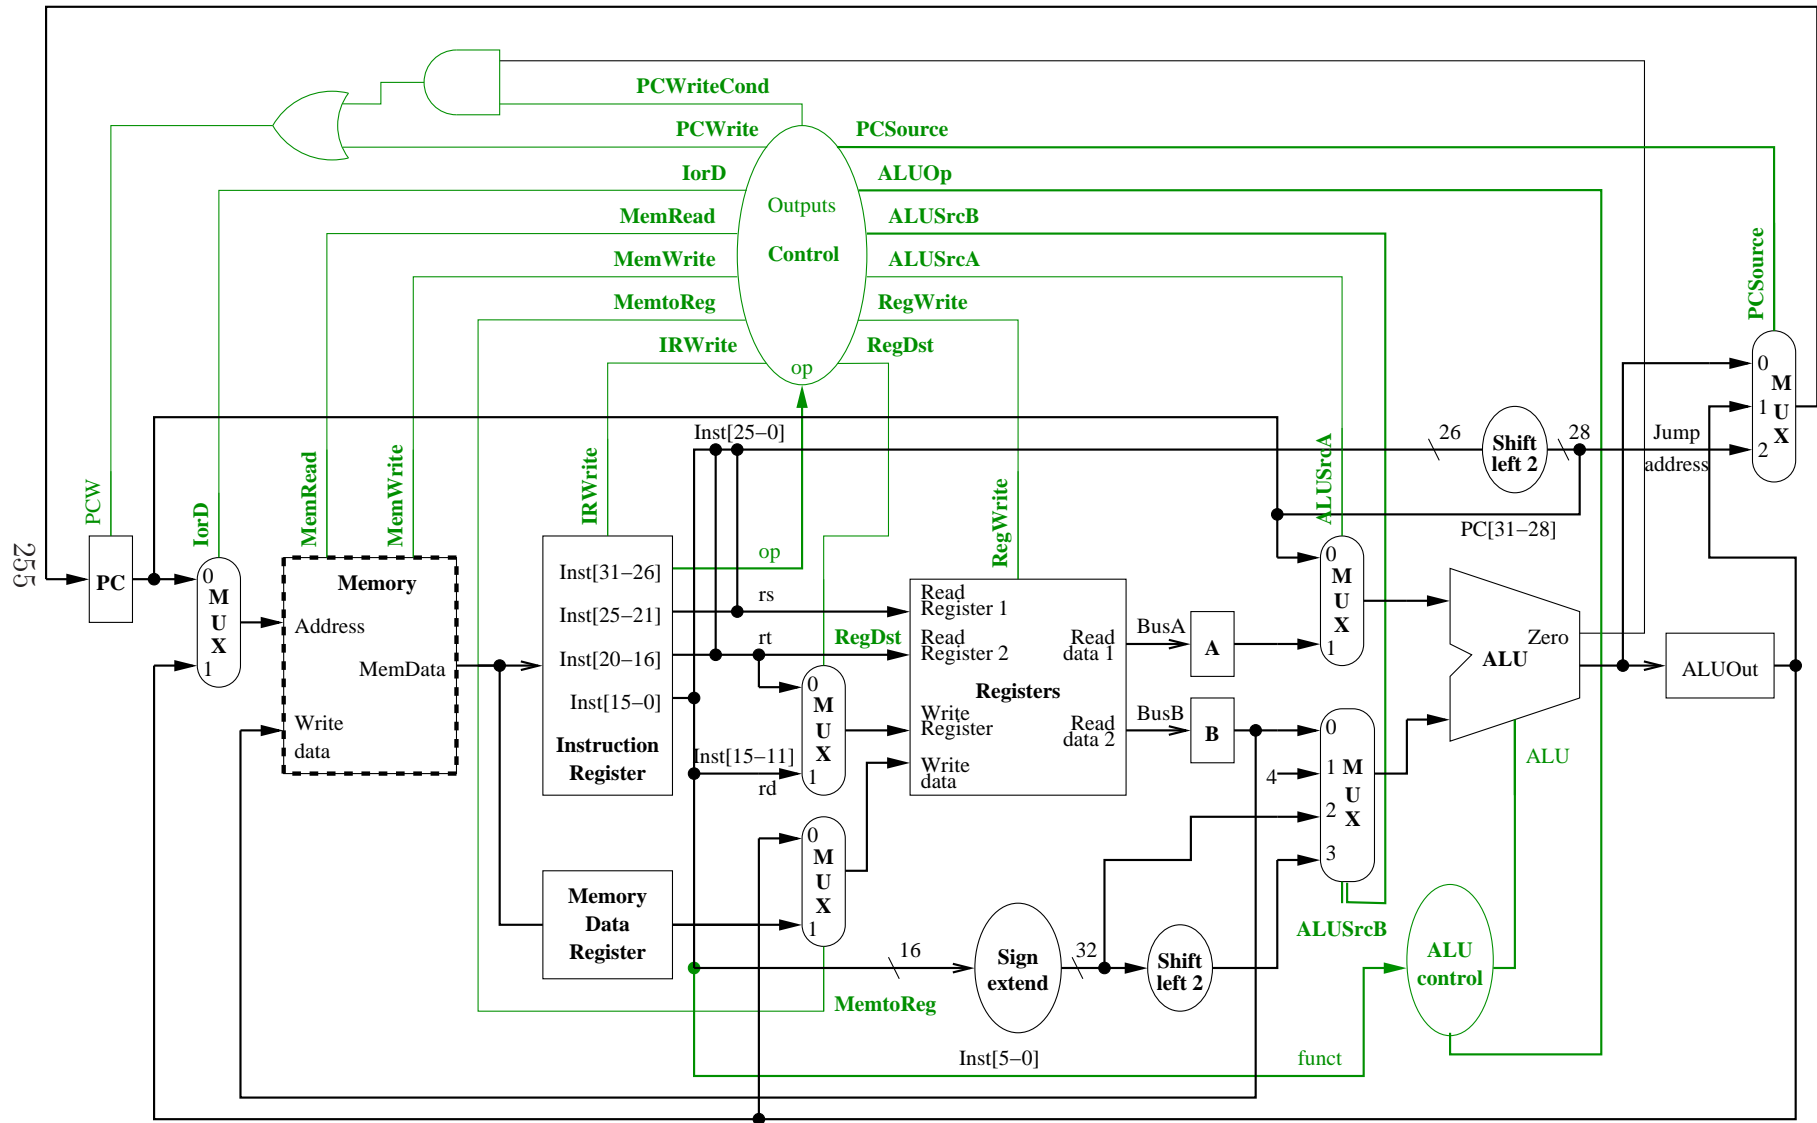












The control signals

The following control signals are identified in the datapath:

Signal	Action when	
	0 (deasserted)	1 (asserted)
RegDst	the register written is the rt field	the register written is the rd field
RegWrite	the register file will not be written into	the register addressed by the instruction will be written into
ALUSrcA	the first ALU operand is the PC	the first ALU operand is register A
MemRead	no memory read occurs	the contents of memory at the specified address is placed on the data bus
MemWrite	no memory write occurs	the contents of register B is written to memory at the specified address
MemtoReg	the value written to the register file comes from ALUOut	the value written to the register file comes from the MDR

Signal	Action when	
	0 (deasserted)	1 (asserted)
IorD	the memory address comes from the PC (an instruction)	the memory address comes from ALUOut (a data read)
IRWrite	the IR is not written into	the IR is written into (an instruction is read)
PCWrite	none (see below)	the PC is written into; the value comes from the MUX controlled by the signal PCSource
PCWriteCond	if both it and PCWrite are not asserted, the PC is not written	the PC is written if the ALU output Zero is active

Following are the 2-bit control signals:

Signal	Value	Action taken
ALUOp	00	ALU performs ADD operation
	01	ALU performs SUBTRACT operation
	10	ALU performs operation specified by funct field
ALUSrcB	00	the second ALU operand is from register B
	01	the second ALU operand is 4
	10	the second ALU operand is the sign extended low order 16 bits of the IR (imm16)
	11	the second ALU operand is the sign extended low order 16 bits of the IR shifted left by 2 bits
PCSource	00	the PC is updated with the value $PC + 4$
	01	the PC is updated with the value in register ALUOut (the branch target address, for a branch instruction)
	10	the PC is updated with the jump target address

The control unit must now be designed.

Since the instructions will now require several states, the control will be a state machine, with the instruction op codes as inputs and the control signals as outputs.

Review of instruction cycles and actions

Cycle	Instruction type	action
IF	all	$IR \leftarrow \text{Memory}[PC]$ $PC \leftarrow PC + 4$
ID	all	$A \leftarrow \text{Reg}[rs]$ $B \leftarrow \text{Reg}[rt]$ $ALUOut \leftarrow PC + (\text{imm16} \ll 2)$
EX	R-type Load/Store Branch Jump	$ALUOut \leftarrow A \text{ op } B$ $ALUOut \leftarrow A + \text{sign-extend}(\text{imm16})$ if $(A == B)$ then $PC \leftarrow ALUOut$ $PC \leftarrow PC[31:28] \parallel (IR[25:0] \ll 2)$
MEM	Load Store	$MDR \leftarrow \text{Memory}[ALUOut]$ $\text{Memory}[ALUOut] \leftarrow B$
WB	R-type Load	$\text{Reg}[rd] \leftarrow ALUOut$ $\text{Reg}[rt] \leftarrow MDR$

Note that the first two steps are required for all instructions, and all instructions require at least the first 3 cycles.

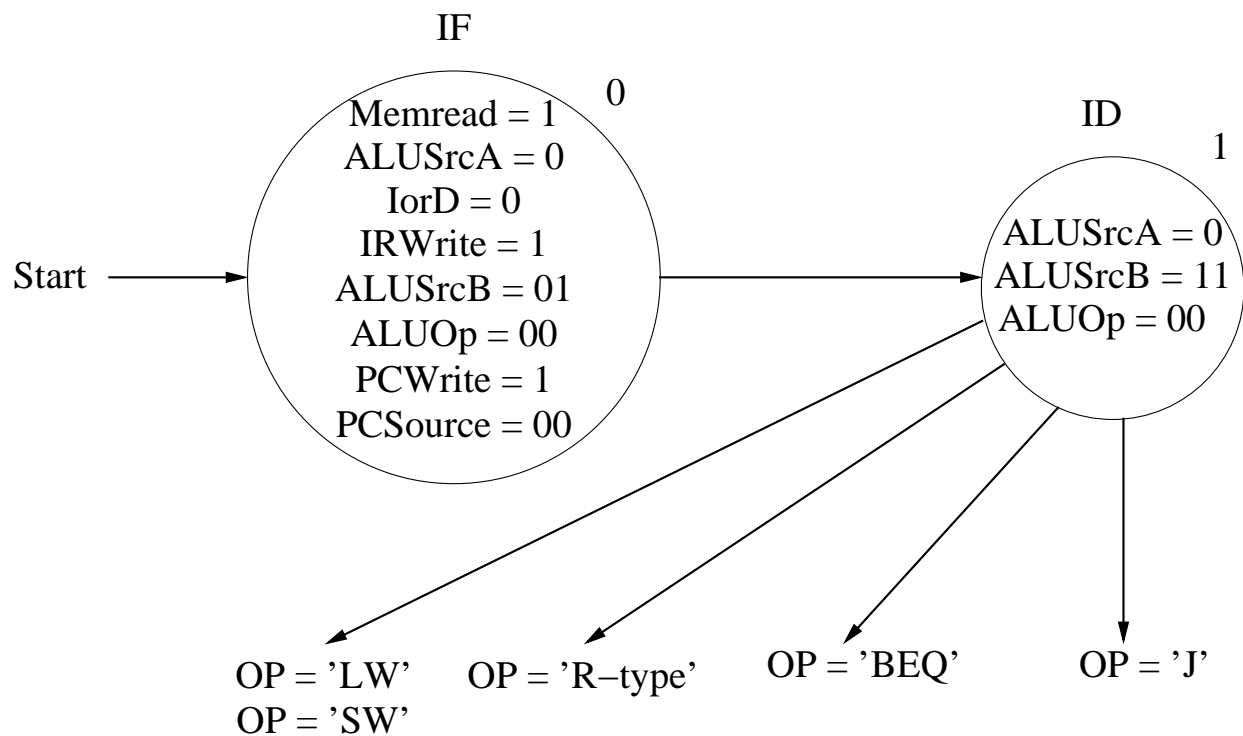
The MEM step is required only by the load and store instructions.

The ALU control unit is still a combinational logic block, as before.

Design of the control unit

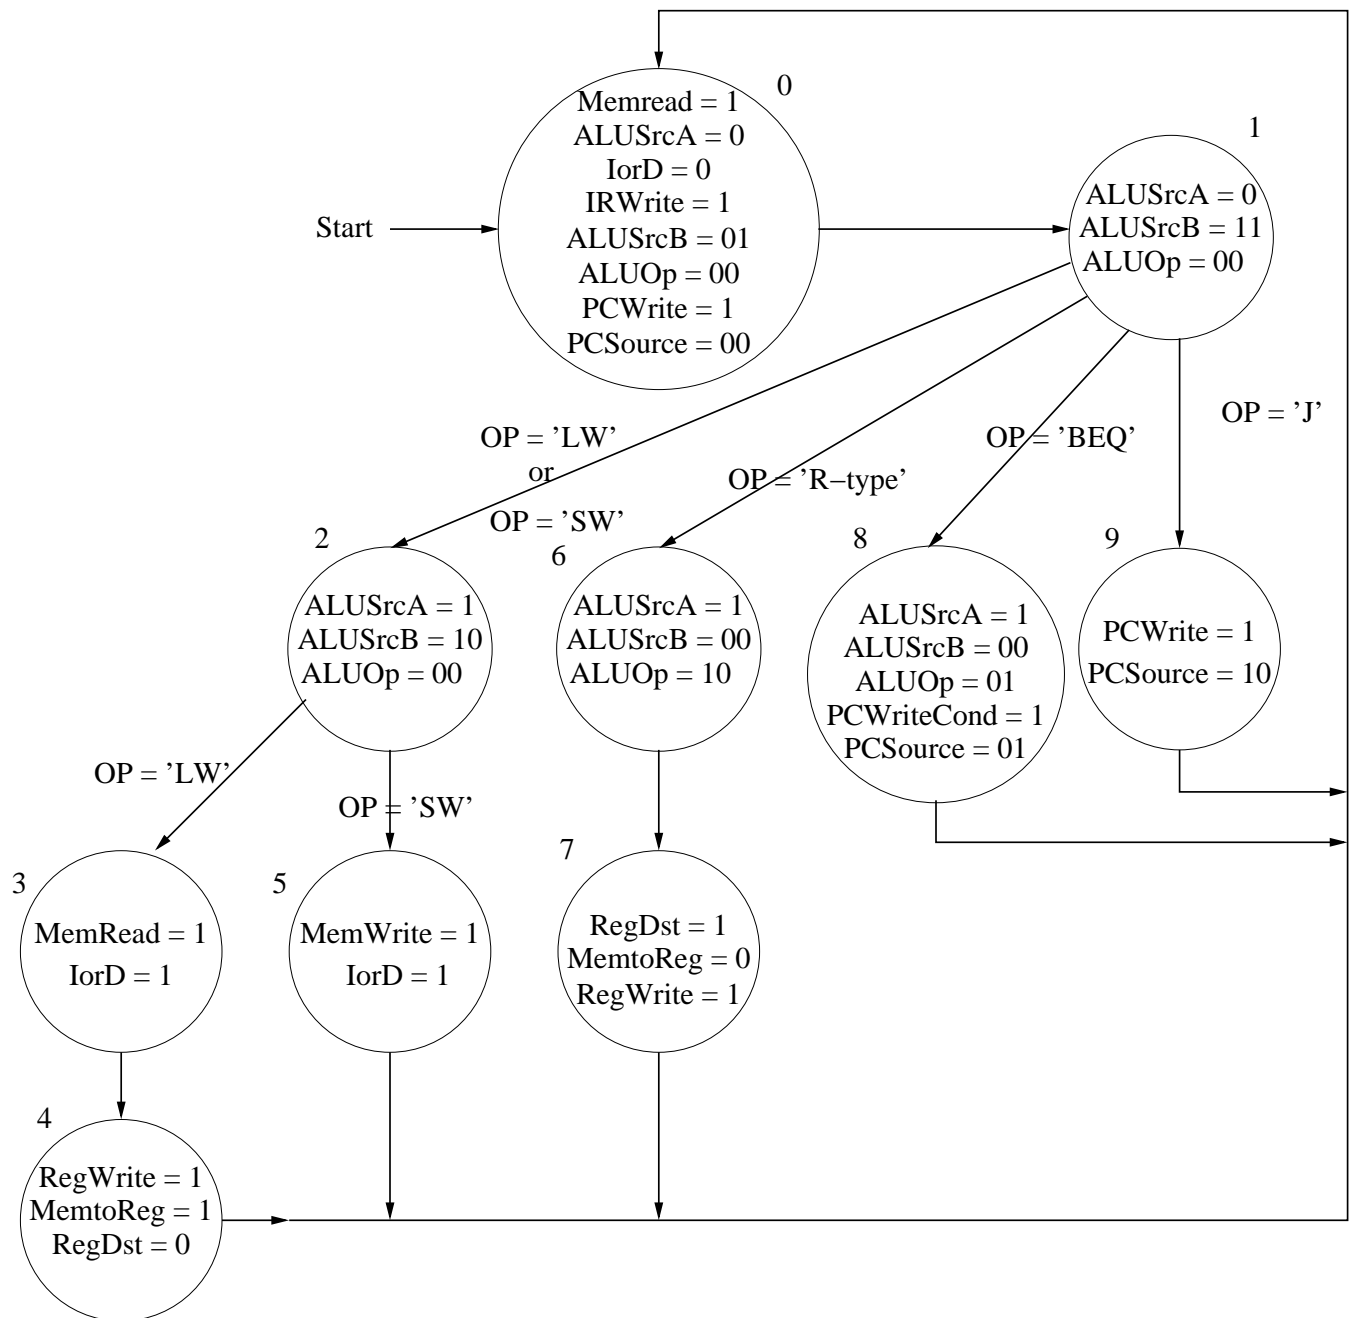
The control unit is a state machine, implementing the state sequencing for every instruction.

Following is a partial state machine, detailing the IF and ID stages, which are the same for all instructions:



The partial state machines which implement each of the instructions follow.

The combined control unit



Implementing the control unit

All that remains is to implement the control unit is to design the control logic itself.

Inputs are the instruction op codes, as before, and the outputs are the control signals.

The following steps are typically followed in the implementation of any sequential device:

- Construct the state diagram or equivalent (done).
- Assign numeric (binary) values to the states.
- Choose a memory element for state memory. (Normally, these would be D flip flops or JK flip flops.)
- Design the combinational logic blocks to implement the next-state functions.
- Design the combinational logic blocks to implement the outputs.

The actual implementation can be done in a number of ways; as discrete logic, a PLA, read-only memory, etc.

Typically, the control unit would be automatically generated from a description in some high level design language.

Adding exception handling

We will implement the hardware and control functions to handle two types of exceptions; *undefined instruction* and *arithmetic overflow*. Recall that the ALU had an overflow detection output, which can be used as an input to the controller.

1. We will use a register labeled **Cause** to store a number (0 or 1) to identify the type of exception, (0 for undefined instruction, 1 for arithmetic overflow).

It requires a control signal **CauseWrite** to be generated by the controller. The controller also must set the value written to the register, depending on whether or not the exception was an arithmetic overflow.

The control signal **IntCause** is used to set this value.

2. The PC will be set to memory address **C0000000** where the operating system is expected to provide an event handler.

This is accomplished by adding another input (input 3) to the MUX which updates the PC address. The MUX is controlled by the 2-bit signal **PCSource**.

3. The address of the instruction which caused the exception is stored in the register **EPC**, a 32 bit register.

Writing to this register is controlled by the new signal **EPCWrite**.

Storing the address of the instruction can be done several ways; for example, it could be stored at the beginning of each instruction.

This would require a change to the datapath, *and* a way to disable the storing of the address after each exception.

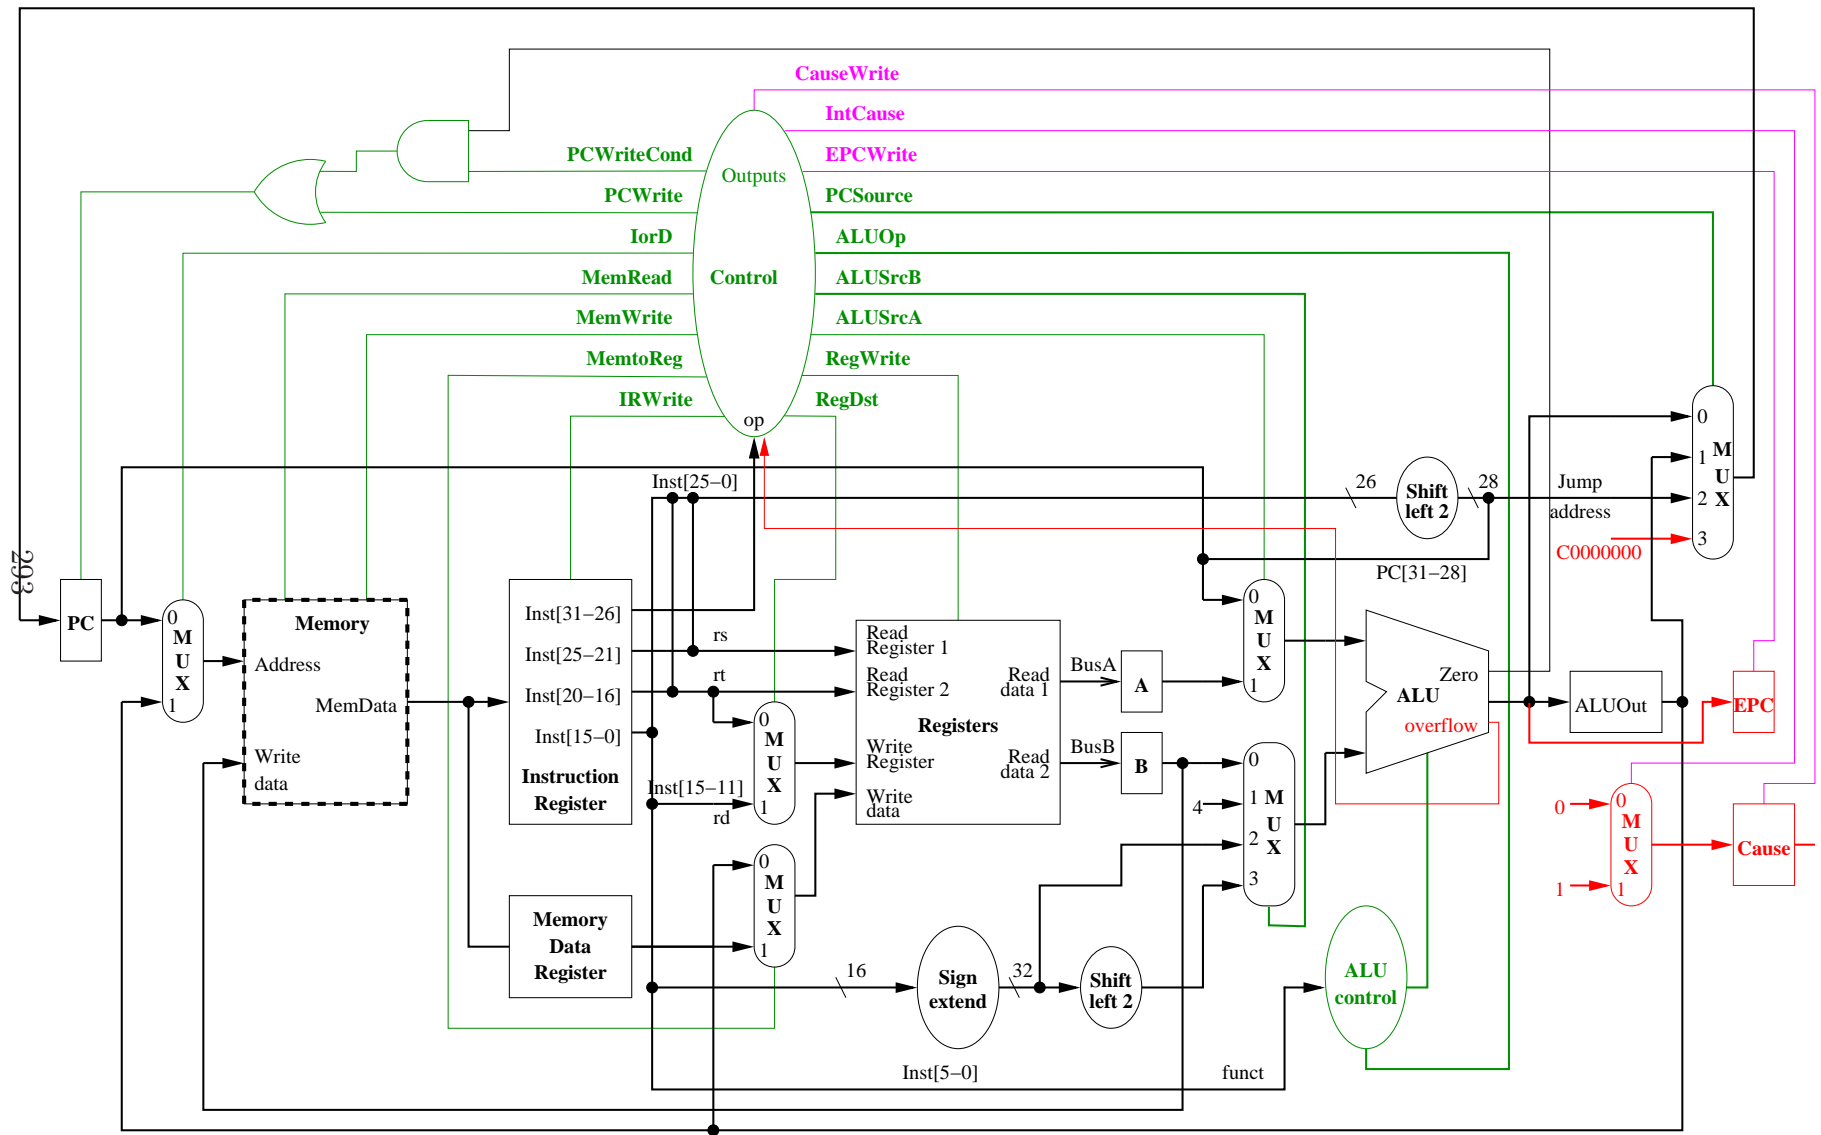
It is possible to store the address with only a small change to the datapath (merely adding the **EPC** register to accept the output of the ALU).

Recall that the next address (**PC** + 4) is calculated in the ALU, and is written to the PC in the first cycle of every instruction. The ALU can be used to subtract the value 4 from the PC after an exception is detected, but before it is written into the EPC, so it contains the actual address of the present instruction.

(Actually, there would be no real problem with saving the value **PC** + 4 in the EPC; the interrupt handler could be responsible for the subtraction.)

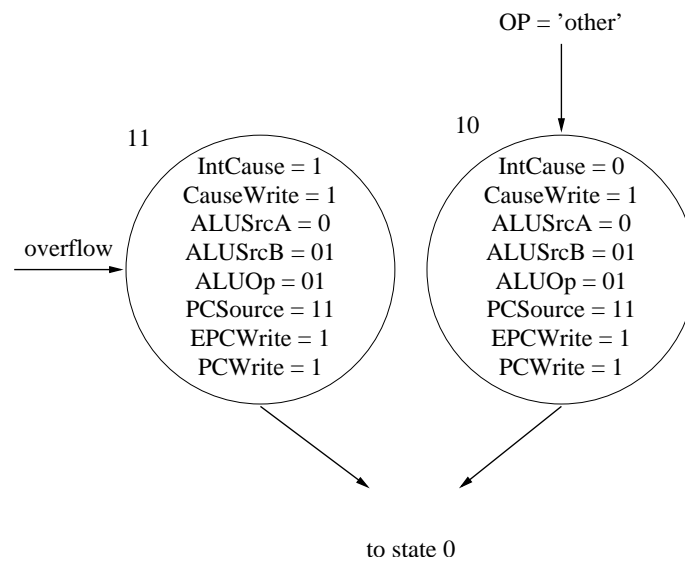
So, in order to handle these two exceptions, we have added two registers — **EPC** and **Cause**, and three control signals — **EPCWrite**, **IntCause**, and **CauseWrite**.

The changes to the processor datapath and control signals required for the implementation of the exceptions detailed above are shown in the following diagram.



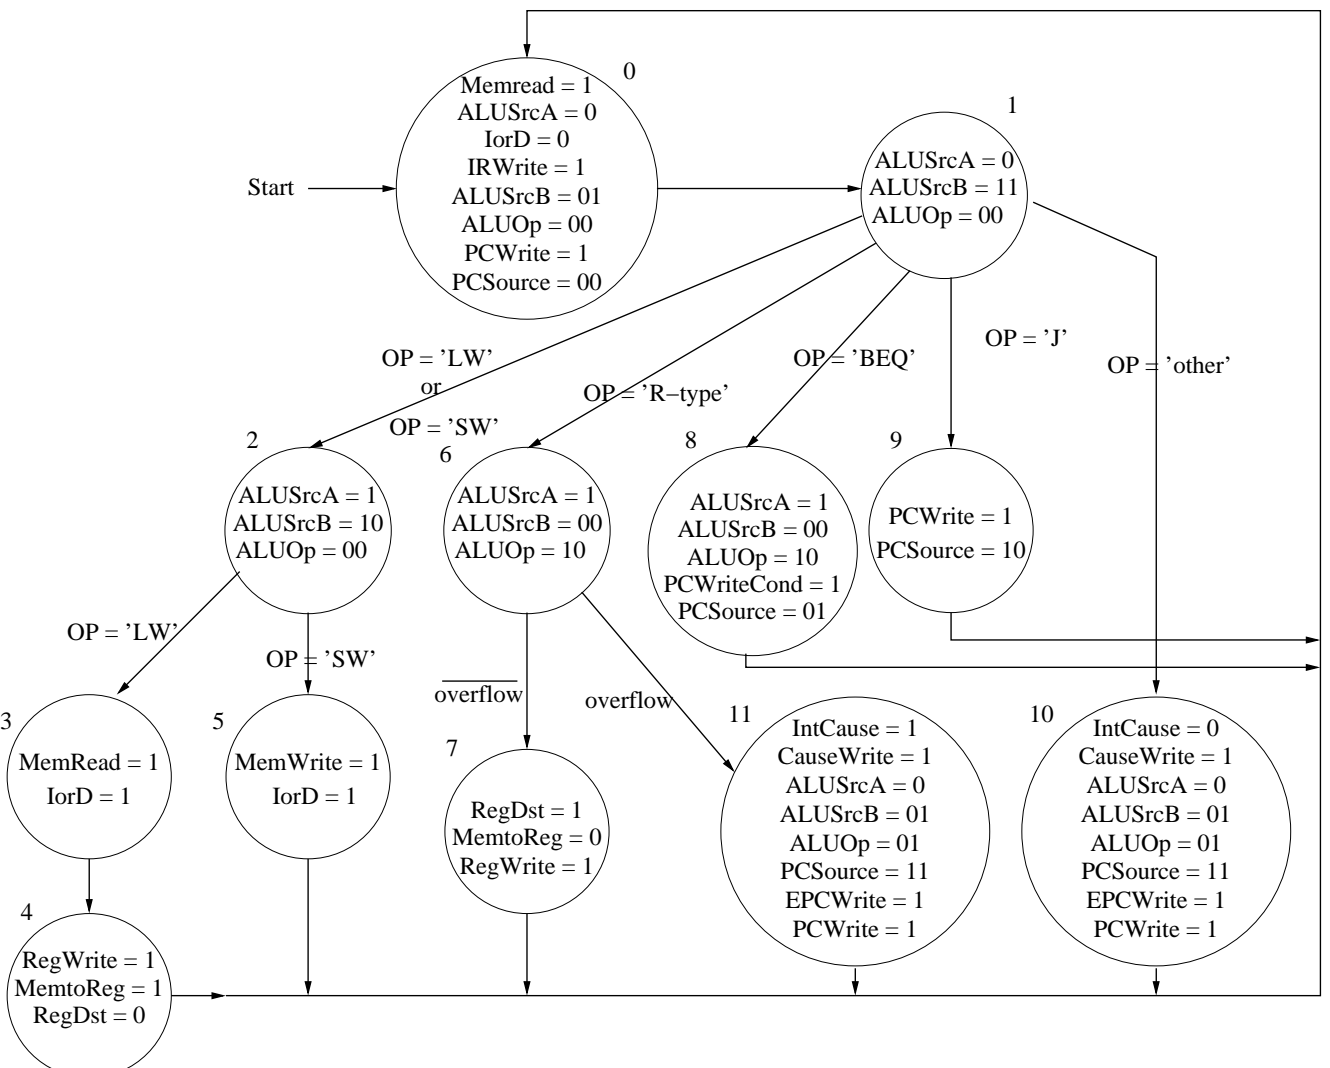
Adding exception handling to the control unit

The exceptions **overflow** and **undefined** can be implemented by the addition of only one state each:



The input **overflow** is an output from the ALU. It is a combinational logic output, produced while the ALU is performing the selected operation.

The control unit, with exception handling



The ALU operation which could result in an overflow is done in the **EX** cycle, and the **overflow** signal is *only* available then, unless it is saved in a register.

More about interrupts

The ability to handle interrupts and exceptions is an important feature for processors.

We have added the control logic to detect the two types of exceptions described earlier, but note that the **Cause** and the **EPC** register cannot be read.

Instructions would have to be provided to allow these registers to be read and manipulated.

Processors usually have *policies* relating to exceptions. The MIPS processor had the policy that instructions which cause an exception has no effect (e.g., nothing is written into a register.)

For some exceptions, if this policy is used, the operation may have to complete before the exception can be detected, and the result of the operation must then be “rolled back.”

This makes the implementation of exceptions difficult — sometimes the state prior to an operation must be saved so it can be restored. This constraint alone sometimes results in instructions requiring more cycles for their implementation.

Exceptions and interrupts in other processors

A common type of interrupt is a **vectored interrupt**. Here, different interrupts or exceptions jump to different addresses in memory. The operating system places an appropriate interrupt handler for the particular interrupt at each of these locations.

A vectored interrupt both identifies the type of interrupt, and provides the handler at the same time. (Since different interrupts or exceptions have different vectors.)

In the INTEL processors, it is the responsibility of the interrupting device to provide the interrupt vector. (This is usually done by one of the peripheral controller chips, under control of the operating system.)

A major problem with the PC architecture is that only a small number of interrupts (typically 16) can be handled by the controller chip. this has lead to many problems with hardware devices “sharing interrupts” — defeating the advantages of vectored interrupts.

We will look at interrupts again, later, when we discuss input and output devices.

Some questions about exceptions and interrupts

The following questions often have different answers for different processors:

- How does a processor return control of the program flow from the exception or interrupt handler to the interrupted program?

Some processors have explicit instructions for this (e.g., the MIPS processors), others treat interrupts and exceptions as being similar to subprogram calls (INTEL processors do this.)

- What happens when an exception or interrupt is itself interrupted?

Some processors save the return addresses in a stack data structure, and successive levels of interrupts just increase the stack depth. Typically, this is the way subprogram return addresses are also stored.

Some processors automatically turn off the interrupt capability at the beginning of an interrupt, and it must be explicitly turned back on by the interrupt or exception handler to accept another interrupt.

Some processors have both features — instructions can turn the interrupt capability on and off, and can allow interrupts to be interrupted themselves. (This turns out to be important for implementing certain operating system functions.)

Comments on our implementation of exceptions

Note that our implementation has only one register for the address of the interrupting instruction, and no way to read that address and modify it to resume the program where the exception occurred.

What changes would be required to the instruction set accomplish this?

The simplest solution would probably be to allow only one interrupt at a time, by disabling the interrupt capability, and to provide:

1. An instruction to store the **EPC** in the register file.
2. An instruction to store the **Cause** register in the register file.
3. An instruction to turn on interrupt capability after the *next* instruction completed execution. (This assumes that the next instruction restores the PC to the address of the instruction following the one that caused the exception.)

Note that these would require changes to the datapath and control.

This example was just to give the flavor of the problems involved with handling exceptions in the processor. More complex instruction sets and architectures exacerbate the problems.

Comments on handling interrupts

Although exception handling is complex, it is often simpler than the handling of external interrupts.

Exceptions occur as a result of occurrences *internal* to the processor. Consequently, they are usually both predictable, and occur and are detected at known times in the execution of a particular instruction.

Interrupts are *external* events, and are *not* at all synchronized with the execution of instructions in the processor.

Since interrupts may be notification of an urgent event, they usually require fast servicing.

Decisions therefore have to be taken about *exactly when* in the execution of an instruction an interrupt will be detected and handled.

Some of the considerations are:

- If the instruction is not allowed to complete, information must be retained in order to either continue or restart the interrupted instruction. How will this be done?
- If the interrupted instruction is allowed to complete, how will the processor return to the next instruction in the current program?
- Can the interrupt handler be interrupted?
- Can interrupts be prioritized so that a high priority interrupt can interrupt a lower priority interrupt?